

Supplementary Material for Variance-Corrected Multi-Asset Equity Simulation with Hybrid Hidden Markov Marginals

S1. Empirical cross-term diagnostic

The variance correction of Eq. (8) assumes $\text{Cov}(\tilde{g}_i^c, g_m) \approx 0$, allowing the variance derivation to omit the cross-term. This assumption does not follow automatically from the fitting procedure because each JumpHMM uses the asset’s full growth-rate history, including its market-correlated component. A remaining market loading in the generator draw would invalidate the closed-form scale. We tested the assumption by sampling $R = 100$ generator growth-rate paths $\tilde{g}_i^{(r)}$ per ticker from each generator and computing the normalized cross-covariance against the real return path for a broad-market exchange-traded fund (ticker SPY):

$$\text{cov}_i^{\text{norm}} = \frac{1}{R} \sum_{r=1}^R \frac{\text{Cov}((\tilde{g}_i^{(r)})^c, g_m)}{\sigma_{\text{gen},i} \sigma_m}. \quad (\text{S1})$$

Centering does not change covariance, so these values are numerically identical to those obtained from the raw draws. We plotted the per-ticker values as histograms (Supplementary Fig. S1). Across the 423 tickers, the full-return JumpHMM generator (used by the hybrid and naive composition methods), the residual-fit JumpHMM generator (used by the JumpHMM-on-residuals composition method), and the Gaussian baseline had universe-mean normalized cross-covariances of 1.53×10^{-4} , 6.58×10^{-5} , and 2.26×10^{-4} , respectively. Their per-ticker values ranged from -0.0061 to 0.0061 . These values quantify the cross-term omitted in Eq. (8).

S2. Rebalancing under independent residuals

To determine what a portfolio-level use of these paths inherits from the independent-residual construction, this appendix derives the diversification return of a rebalanced portfolio under Eq. (4). The variance correction in Eq. (8) sets each synthetic asset’s second moment to its generator target. When that target matches the empirical second moment, the second-order formula for the diversification return is matched at the asset level as well. What the construction does not match is the cross-asset residual covariance, and the derivation below isolates where that omission enters.




Figure S1: **Empirical normalized cross-covariance** $\text{Cov}(\tilde{g}_i^c, g_m)/(\sigma_{\text{gen},i} \sigma_m)$ **across the 423-ticker universe.** Histograms were stacked per generator, with each per-ticker value averaged over $R = 100$ generator growth-rate paths. Universe means were 1.53×10^{-4} , 6.58×10^{-5} , and 2.26×10^{-4} for the full-return JumpHMM, residual-fit JumpHMM, and Gaussian generators.

The synthetic paths reproduce per-asset marginals and single-factor cross-sectional dependence. Each residual generator retains its own temporal structure, but the residual paths are generated independently across assets at each date. The construction therefore omits cross-asset residual covariance. Consider a daily-rebalanced portfolio of N assets with constant target weights $w_i \geq 0$, $\sum_i w_i = 1$, reset to those weights at the start of each trading day. Throughout this derivation $g_i(t)$ is the annualized log-growth rate of Eq. (1), so the per-period log return on asset i is $g_i(t) \Delta t$ with $\Delta t = 1/252$ yr. The wealth process is:

$$W_t = W_{t-1} \sum_i w_i \exp(g_i(t) \Delta t), \quad (\text{S2})$$

so the per-step log return on the portfolio is $\log(W_t/W_{t-1}) = \log \sum_i w_i \exp(g_i(t) \Delta t)$. We suppress the time argument, writing g_i for $g_i(t)$. A second-order Taylor expansion of the exponential about zero is $\exp(g_i \Delta t) = 1 + g_i \Delta t + \frac{1}{2}(g_i \Delta t)^2 + O((g \Delta t)^3)$. Substituting it into each summand of $\sum_i w_i \exp(g_i \Delta t)$ gives:

$$\begin{aligned} \sum_i w_i \exp(g_i \Delta t) &= \sum_i w_i [1 + g_i \Delta t + \frac{1}{2}(g_i \Delta t)^2 + O((g \Delta t)^3)] && (\text{Taylor}) \\ &= \sum_i w_i + \Delta t \sum_i w_i g_i + \frac{1}{2}(\Delta t)^2 \sum_i w_i g_i^2 + O((g \Delta t)^3) && (\text{distribute}) \\ &= 1 + \Delta t \sum_i w_i g_i + \frac{1}{2}(\Delta t)^2 \sum_i w_i g_i^2 + O((g \Delta t)^3). && (\sum_i w_i = 1) \end{aligned} \quad (\text{S3})$$

A second-order Taylor expansion of $\log(1+x)$ about $x=0$ gives $\log(1+x) = x - \frac{1}{2}x^2 + O(x^3)$. Identifying $1+x$ with the right-hand side of Eq. (S3), the

order- Δt correction is $x := \Delta t \sum_i w_i g_i + \frac{1}{2}(\Delta t)^2 \sum_i w_i g_i^2$, and squaring keeps only the leading $(\Delta t)^2$ term: $x^2 = (\Delta t)^2 (\sum_i w_i g_i)^2 + O((\Delta t)^3)$ because the cross term has an extra Δt factor. Substitution gives:

$$\begin{aligned}
\log \frac{W_t}{W_{t-1}} &= \log \left(1 + x + O((g \Delta t)^3) \right) \\
&= x - \frac{1}{2} x^2 + O(x^3) \\
&= \Delta t \sum_i w_i g_i + \frac{1}{2} (\Delta t)^2 \sum_i w_i g_i^2 \\
&\quad - \frac{1}{2} (\Delta t)^2 \left(\sum_i w_i g_i \right)^2 + O((g \Delta t)^3) \\
&= \Delta t \sum_i w_i g_i + \frac{1}{2} (\Delta t)^2 \left[\sum_i w_i g_i^2 - \left(\sum_i w_i g_i \right)^2 \right] + O((g \Delta t)^3). \tag{S4}
\end{aligned}$$

Dividing through by Δt extracts the portfolio's annualized one-step log-growth rate $g_p(t) := \Delta t^{-1} \log(W_t/W_{t-1})$, the portfolio counterpart of the per-asset $g_i(t)$ (still suppressing the time argument):

$$g_p = \sum_i w_i g_i + \frac{1}{2} \Delta t \left(\sum_i w_i g_i^2 - \left(\sum_i w_i g_i \right)^2 \right) + O((\Delta t)^2). \tag{S5}$$

To obtain expected annualized portfolio growth, we restored the time argument and took unconditional expectations of Eq. (S5) under stationarity. Define the per-asset annualized mean and variance:

$$\mu_i := \mathbb{E}(g_i(t)), \quad \sigma_i^2 := \text{Var}(g_i(t)), \tag{S6}$$

the per-rate covariance $\Sigma_{ij} := \text{Cov}(g_i(t), g_j(t))$, and the per-rate portfolio variance $\sigma_p^2 := \text{Var}(\sum_i w_i g_i(t)) = w^\top \Sigma w$. The $O(\Delta t)$ correction in Eq. (S5) contains two squared quantities, $\sum_i w_i g_i(t)^2$ and $(\sum_i w_i g_i(t))^2$. Both expand by the second-moment identity, which holds for any random variable X with finite second moment:

$$\mathbb{E}(X^2) = \mathbb{E}(X)^2 + \text{Var}(X). \tag{S7}$$

Applying this identity to $X = g_i(t)$ and summing against w_i gives:

$$\begin{aligned}
\mathbb{E} \left(\sum_i w_i g_i(t)^2 \right) &= \sum_i w_i \mathbb{E}(g_i(t)^2) && \text{(linearity)} \\
&= \sum_i w_i [\mu_i^2 + \sigma_i^2] && \text{(by Eq. (S7))} \\
&= \sum_i w_i \mu_i^2 + \sum_i w_i \sigma_i^2. \tag{S8}
\end{aligned}$$

Applying the identity to $X = \sum_i w_i g_i(t)$, whose mean is $\sum_i w_i \mu_i$ and whose variance is σ_p^2 , gives:

$$\begin{aligned}\mathbb{E}\left(\left(\sum_i w_i g_i(t)\right)^2\right) &= \mathbb{E}\left(\sum_i w_i g_i(t)\right)^2 + \text{Var}\left(\sum_i w_i g_i(t)\right) \quad (\text{by Eq. (S7)}) \\ &= \left(\sum_i w_i \mu_i\right)^2 + \sigma_p^2.\end{aligned}\tag{S9}$$

We substituted Eqs. (S8)–(S9) into the expectation of Eq. (S5) and grouped the variance and mean contributions as follows:

$$\begin{aligned}\mathbb{E}(g_p) &= \sum_i w_i \mu_i + \underbrace{\frac{1}{2} \Delta t \left(\sum_i w_i \sigma_i^2 - \sigma_p^2 \right)}_D \\ &\quad + \frac{1}{2} \Delta t \left[\sum_i w_i \mu_i^2 - \left(\sum_i w_i \mu_i \right)^2 \right] + O((\Delta t)^2),\end{aligned}\tag{S10}$$

the Booth-Fama identity for the daily-rebalanced portfolio in annualized growth-rate units (Booth and Fama, 1992; Markowitz, 1976). The middle term, D , is the diversification return: the $O(\Delta t)$ correction equal to half the difference between the weighted-average per-asset variance $\sum_i w_i \sigma_i^2$ and the portfolio variance σ_p^2 . The per-rate moments μ_i and σ_i^2 have units yr^{-1} and yr^{-2} respectively; the products $\Delta t \sigma_i^2$ and $\Delta t \sigma_p^2$ in Eq. (S10) are the annualized variances (yr^{-1}). The third term in Eq. (S10) is the cross-sectional dispersion of the per-asset annualized means, of order $\Delta t \mu^2 \sim 3 \times 10^{-5} \text{ yr}^{-1}$ at $\mu \sim 0.08 \text{ yr}^{-1}$ and negligible against the variance term $D \sim \Delta t \sigma^2 \sim 4 \times 10^{-2} \text{ yr}^{-1}$; we drop it. The realized annualized compound growth rate $\bar{g}_p = (T \Delta t)^{-1} \log(W_T/W_0) = T^{-1} \sum_t g_p(t)$ converges to $\mathbb{E}(g_p)$ by the ergodic theorem under any stationary residual process, so the long-run limit is fixed by the one-time-slice unconditional joint distribution of $(g_i(t))_i$ and does not, in itself, depend on the serial structure of the residual draw.

To rewrite D in a form that makes its non-negativity manifest, expand $\sigma_p^2 = \sum_i w_i^2 \sigma_i^2 + 2 \sum_{i < j} w_i w_j \Sigma_{ij}$ and use the identity $\sum_i w_i (1 - w_i) \sigma_i^2 = \sum_{i < j} w_i w_j (\sigma_i^2 + \sigma_j^2)$ that holds for any weights summing to one:

$$D = \frac{1}{2} \Delta t \left(\sum_i w_i \sigma_i^2 - \sigma_p^2 \right) = \frac{1}{2} \Delta t \sum_{i < j} w_i w_j \text{Var}(g_i - g_j) \geq 0, \tag{S11}$$

so D equals one-half the annualized weighted-average pairwise growth-rate-difference variance and is therefore non-negative, with equality only when every pair (i, j) comoves perfectly ($\text{Var}(g_i - g_j) = 0$) or all weight concentrates on a single asset. The rebalancing rule itself produces this effect: weights drift between rebalances, the rebalancer sells whatever has outperformed (its weight rose above target) and buys whatever has underperformed (its weight fell below target), and the trade systematically converts period-by-period cross-sectional

dispersion into compounded growth. Markowitz (1976) described this as the variance bonus of the long-run mean-variance allocation; Booth and Fama (1992) called it the diversification return.

To separate market and residual contributions to the diversification return, we substituted the single-index decomposition $g_i(t) = \alpha_i + \beta_i g_m(t) + \varepsilon_i(t)$, with market variance $\sigma_m^2 := \text{Var}(g_m(t))$, per-asset residual variance $\sigma_{\varepsilon,i}^2 := \text{Var}(\varepsilon_i(t))$, and $\text{Cov}(g_m, \varepsilon_i) = 0$ (verified in Sec. S1), into Eq. (S11) gives:

$$\text{Var}(g_i - g_j) = (\beta_i - \beta_j)^2 \sigma_m^2 + \text{Var}(\varepsilon_i - \varepsilon_j), \quad (\text{S12})$$

Because g_m and the residuals are uncorrelated, this relation splits the diversification return into market-beta-dispersion and idiosyncratic components:

$$D = \underbrace{\frac{1}{2} \Delta t \sigma_m^2 \sum_{i < j} w_i w_j (\beta_i - \beta_j)^2}_{D_{\text{mkt}}} + \underbrace{\frac{1}{2} \Delta t \sum_{i < j} w_i w_j \text{Var}(\varepsilon_i - \varepsilon_j)}_{D_{\varepsilon}}, \quad (\text{S13})$$

where the market-beta-dispersion component D_{mkt} equals $\frac{1}{2} \Delta t \sigma_m^2 [\sum_i w_i \beta_i^2 - (\sum_i w_i \beta_i)^2]$, the weighted cross-sectional β dispersion scaled by the annualized market variance $\Delta t \sigma_m^2$. Across the 423-ticker universe the empirical β distribution spans $[0.03, 1.79]$ with interquartile range $[0.81, 1.18]$ and the SPY annualized variance $\Delta t \sigma_m^2 \approx 0.04 \text{ yr}^{-1}$. The weighted variance of β is bounded by one quarter of its squared range, so D_{mkt} is at most 0.016 yr^{-1} , or 1.6 percentage points per year over this universe. For the idiosyncratic component D_{ε} , residuals were approximately orthogonal across assets at the same time slice. Therefore, $\text{Var}(\varepsilon_i - \varepsilon_j) \approx \sigma_{\varepsilon,i}^2 + \sigma_{\varepsilon,j}^2$, and D_{ε} scales with the weighted-average annualized idiosyncratic variance $\Delta t \sigma_{\varepsilon,i}^2$ per name.

The second-order value of D depends on same-time unconditional moments, not on serial dependence. Equation (8) sets $\sigma_{\text{gen},i}^2 = \beta_i^2 \sigma_m^2 + \sigma_{\varepsilon,i}^2$ to its empirical target. Any residual generator with the same unconditional second moments therefore gives the same D . This includes a stationary first-order autoregressive [AR(1)] residual generator $\varepsilon_i(t) = \phi \varepsilon_i(t-1) + \eta_i(t)$ with independent and identically distributed innovations η_i tuned so that $\text{Var}(\varepsilon_i) = \sigma_{\eta,i}^2 / (1 - \phi^2) = \sigma_{\varepsilon,i}^2$. The parameter ϕ leaves the marginal distribution at a single time slice unchanged. It therefore does not change $\mathbb{E}(g_p)$ in Eqs. (S10) and (S13), but it changes sampling variance. For a stationary AR(1) residual with variance σ_{ε}^2 , the variance of its sample mean is approximately $T^{-1} \sigma_{\varepsilon}^2 (1 + \phi) / (1 - \phi)$. Under ergodicity, serial dependence changes per-path sampling noise, not the stationary mean. Serial dependence in the residual draw is therefore not a lever on the diversification return at the level of the second-order formula.

The direct structural difference is contemporaneous. The derivation assumes residuals are orthogonal across assets at each time slice, as they are in the independent per-asset generator. Real residuals share sector and style factors, which reduce $\text{Var}(\varepsilon_i - \varepsilon_j)$ below the independent-generator value $\sigma_{\varepsilon,i}^2 + \sigma_{\varepsilon,j}^2$. Matching each marginal variance does not restore this covariance. The synthetic D_{ε} will therefore exceed its real-market counterpart when the omitted covariance

is positive. Quantifying that contribution requires a joint-residual experiment with the marginals and the allocation rule held fixed, which we did not run. Realized diversification returns on real equity portfolios are also attenuated relative to the second-order formula at all rebalancing horizons (Bouchey et al., 2015), so the synthetic excess is an upper bound on what an independent-residual generator would contribute in a live comparison.

S3. Validation metric definitions

To make the validation criteria reproducible, we defined the distributional and tail metrics reported in the main results. The two-sample Kolmogorov–Smirnov (KS) statistic between the empirical cumulative distribution functions (CDFs) of the observed and simulated growth-rate samples (G_{obs} of length n and G_{sim} of length m) is $D = \sup_x |F_n(x) - F_m(x)|$ (Kolmogorov, 1933; Smirnov, 1948); the test rejects when D exceeds the critical value at level $\alpha = 0.05$. The Anderson–Darling (AD) statistic places greater weight on the tails via the integral $A^2 = nm/(n+m) \int [F_n(x) - F_m(x)]^2 / [F(x)(1-F(x))] dF(x)$, where F is the pooled empirical CDF (Anderson and Darling, 1952). The Wasserstein-1 distance between two empirical distributions of equal length is the average of sorted-quantile differences:

$$W_1 = \frac{1}{n} \sum_{i=1}^n |G_{\text{obs}}^{(i)} - G_{\text{sim}}^{(i)}|, \quad (\text{S14})$$

with $G^{(i)}$ the i th order statistic. The Hill upper-tail index is the classical order-statistic estimator:

$$\hat{\xi} = \frac{1}{k-1} \sum_{i=1}^{k-1} \log \left(\frac{X_{(i)}}{X_{(k)}} \right), \quad (\text{S15})$$

with $X_{(1)} \geq X_{(2)} \geq \dots$ the ordered upper-tail absolute returns; k is set to the upper 5% of the sample. For a Pareto-tailed distribution with shape α , $\hat{\xi} \approx 1/\alpha$. For the Hellinger calculation, we binned both samples on a common grid and normalized the bin counts to probability masses p and q . Their distance is:

$$H(p, q) = \frac{1}{\sqrt{2}} \left\| \sqrt{p} - \sqrt{q} \right\|_2 \in [0, 1], \quad (\text{S16})$$

with $H = 0$ for identical probability masses and $H = 1$ for disjoint supports.

We abbreviate mean absolute error in the autocorrelation function as ACF-MAE.

S4. Derivation of the single-asset generator

This appendix derives the probability law of the fitted single-asset generator and connects each fitted component to the empirical feature it controls. The




Figure S2: **Statistical validation of the hidden Markov model with jumps (HMM-WJ)** ($N = 100$, $\alpha = 0.05$). Panels (a) and (d) condition on the $\sim 24\%$ of in-sample paths containing jump events. (a) In-sample KS p -values for 239 jump paths (pass rate 93.7%). (b) Out-of-sample KS p -values for all 1,000 paths (pass rate 95.4%). (c) Out-of-sample density fan chart; the observed density falls within the simulation envelope. (d) Observed and mean simulated in-sample autocorrelation function (ACF) of $|G_t|$ for jump paths.

Table S1: **Descriptive statistics and stylized-facts tests for SPY daily excess growth rates**. In-sample (IS): 2014–2024 ($T = 2,766$); out-of-sample (OoS): 2025 ($T = 249$). Gaussian and Laplace columns show maximum likelihood estimation (MLE) fits to the in-sample data. We used the Jarque–Bera (JB) normality test and the Ljung–Box (LB) autocorrelation test at lag 20. * denotes rejection at $\alpha = 0.05$.

Statistic	IS Observed	OoS Observed	Gaussian	Laplace
Mean (annualized, %)	6.31	11.60	6.31	14.19
Standard deviation of G_t (% yr ⁻¹)	214.50	220.62	214.46	205.93
Skewness	-0.753	-1.093	0	0
Excess Kurtosis	7.715	6.867	0	3
JB normality test	reject* ($p < 0.001$)	reject* ($p < 0.001$)	not used	not used
LB test on G_t (lag 20)	reject* ($p < 0.001$)	reject* ($p = 0.019$)	not used	not used
LB test on $ G_t $ (lag 20)	reject* ($p < 0.001$)	reject* ($p < 0.001$)	not used	not used

hidden Markov model with jumps (HMM-WJ), its no-jump restriction (HMM-NJ), and their symbols follow the main-text definitions in Eq. (2) and Algorithm 1. Unless noted otherwise, numerical values evaluate the closed-form expressions at the fitted SPY model of Table 1: $N = 100$ states, Student- t degrees of freedom $\nu = 5$, jump probability $\epsilon = 10^{-4}$, mean episode duration $\lambda = 90$ steps, tail-selection probability $p_{\text{neg}} = 0.52$, tail-set size $N_{\text{tail}} = 5$, and the in-sample series of $T = 2,766$ daily excess growth rates.

S4.1. Generative law

The no-jump model generates a hidden state path $S_0, \dots, S_T$ on $\{1, \dots, N\}$ and observations $G_1, \dots, G_T$ with the joint density:

$$p(s_0, \dots, s_T, g_1, \dots, g_T) = \bar{\pi}_{s_0} \prod_{t=1}^T \mathbf{T}_{s_{t-1}, s_t} f_{s_t}(g_t), \quad (\text{S17})$$

where $\bar{\pi}_{s_0}$ is the initial probability of state s_0 , $\mathbf{T}_{s_{t-1}, s_t}$ is the probability of moving from state s_{t-1} to state s_t , and f_k is the emission density of state k . Under Eq. (2), f_k is the location-scale Student- t density:

$$f_k(g) = \frac{1}{\sigma_k} \psi_\nu\left(\frac{g - \mu_k}{\sigma_k}\right), \quad \psi_\nu(x) = \frac{\Gamma(\frac{\nu+1}{2})}{\sqrt{\nu\pi} \Gamma(\frac{\nu}{2})} \left(1 + \frac{x^2}{\nu}\right)^{-\frac{\nu+1}{2}}, \quad (\text{S18})$$

where ψ_ν is the density of a standard Student- t random variable with ν degrees of freedom, Γ is the gamma function, and (μ_k, σ_k) are the location and scale of state k . The factorization in Eq. (S17) states the structural property that drives everything that follows: given the state path, the observations are independent. Every form of serial dependence the generator can produce must therefore be carried by the state path (Rabiner and Juang, 1986).

The jump-augmented model enlarges the hidden state. Let $J_t \in \{0, 1, 2, \dots\}$ count the forced steps remaining after time t , and let u be the forced-step state distribution, which mixes uniform draws over the two tail sets:

$$u_k = \begin{cases} p_{\text{neg}}/N_{\text{tail}} & k \in \mathcal{S}_-, \\ (1 - p_{\text{neg}})/N_{\text{tail}} & k \in \mathcal{S}_+, \\ 0 & \text{otherwise.} \end{cases} \quad (\text{S19})$$

Writing $p_K(m) = e^{-\lambda} \lambda^m / m!$ for the Poisson probability of episode length m and $\mathbf{1}\{\cdot\}$ for the indicator function, the pair (S_t, J_t) evolves by the transition kernel:

$$\begin{aligned} & \Pr(S_t = k, J_t = j \mid S_{t-1} = s, J_{t-1} = i) \\ &= \begin{cases} u_k \mathbf{1}\{j = i - 1\} & i \geq 1, \\ (1 - \epsilon + \epsilon e^{-\lambda}) \mathbf{T}_{s,k} \mathbf{1}\{j = 0\} + \epsilon p_K(j + 1) u_k & i = 0, \end{cases} \end{aligned} \quad (\text{S20})$$

and the emission draw depends only on the first coordinate, $G_t \sim f_{S_t}$. The two cases read as follows. While an episode is active ($i \geq 1$), the next state is

an independent draw from u and the counter decrements; the tail set and the state within it are re-drawn at every forced step. While no episode is active ($i = 0$), the chain either takes an ordinary transition (no trigger, or a trigger whose Poisson draw returns zero) or starts an episode of length $j + 1 \geq 1$ whose first forced step occurs immediately. The pair (S_t, J_t) is therefore itself a Markov chain, and the HMM-WJ generator is a hidden Markov model on the enlarged state space $\{1, \dots, N\} \times \{0, 1, 2, \dots\}$. Algorithm 1 implements Eq. (S20) with the episode counter truncated at the simulation horizon; it emits the stationary initial draw as the first observation and applies the kernel from the second step onward, which changes the law of a path only through the absent possibility of a trigger at the first step. When an episode ends, ordinary transitions resume from the last forced tail state, so the chain re-enters the body of the state space through the tail rows of $\mathbf{T}$.

S4.2. Estimators and the one-step empirical law

The fitting procedure consists of four closed-form estimators, and each one matches a specific empirical object exactly. The Laplace partition comes first. For observations $G_1, \dots, G_T$, the Laplace log-likelihood with location μ_L and scale b_L is:

$$\ell(\mu_L, b_L) = -T \ln(2b_L) - \frac{1}{b_L} \sum_{t=1}^T |G_t - \mu_L|. \quad (\text{S21})$$

For any fixed $b_L > 0$, maximizing ℓ over μ_L minimizes $\sum_t |G_t - \mu_L|$, so $\hat{\mu}_L$ is the sample median. Setting the b_L derivative to zero at $\hat{\mu}_L$ gives:

$$\begin{aligned} \frac{\partial \ell}{\partial b_L} &= -\frac{T}{b_L} + \frac{1}{b_L^2} \sum_{t=1}^T |G_t - \hat{\mu}_L| = 0 \quad (\text{first-order condition}) \\ \hat{b}_L &= \frac{1}{T} \sum_{t=1}^T |G_t - \hat{\mu}_L|, \quad (\text{solve for } b_L) \end{aligned} \quad (\text{S22})$$

the mean absolute deviation about the median. The state boundaries then invert the fitted Laplace cumulative distribution function (CDF) (Kotz et al., 2001). The Laplace CDF is piecewise exponential:

$$F_L(x) = \begin{cases} \frac{1}{2} \exp((x - \hat{\mu}_L)/\hat{b}_L) & x \leq \hat{\mu}_L, \\ 1 - \frac{1}{2} \exp(-(x - \hat{\mu}_L)/\hat{b}_L) & x > \hat{\mu}_L, \end{cases} \quad (\text{S23})$$

and setting $F_L(Q_k) = k/N$ and solving each branch for Q_k gives the closed-form boundaries:

$$\begin{aligned} \frac{k}{N} &= \frac{1}{2} e^{(Q_k - \hat{\mu}_L)/\hat{b}_L} \implies Q_k = \hat{\mu}_L + \hat{b}_L \ln \frac{2k}{N}, \quad (\text{lower branch}) \\ \frac{k}{N} &= 1 - \frac{1}{2} e^{-(Q_k - \hat{\mu}_L)/\hat{b}_L} \implies Q_k = \hat{\mu}_L - \hat{b}_L \ln \left(2 - \frac{2k}{N}\right). \quad (\text{upper branch}) \end{aligned} \quad (\text{S24})$$

The bins are equiprobable under the fitted Laplace, not under the empirical distribution, so the in-sample state counts are not uniform: across the 100 states, they range from 16 to 47 of the 2,766 observations. Every within-state mean μ_k and standard deviation σ_k was therefore computed from at least 16 observations, and no state required a fallback to the global moments.

The transition matrix is the maximum likelihood estimator (MLE) of a first-order Markov chain. Conditional on the assigned state sequence $s_1, \dots, s_T$, let n_{jk} count the transitions from state j to state k and let $n_{j\cdot} = \sum_k n_{jk}$ be the row total. The likelihood of the transition parameters factorizes over rows into multinomials, and maximizing the log-likelihood subject to each row summing to one, with multiplier c_j for row j , uses the Lagrangian:

$$\mathcal{L} = \sum_{j,k} n_{jk} \ln \mathbf{T}_{jk} + \sum_j c_j \left(1 - \sum_k \mathbf{T}_{jk}\right), \quad (\text{S25})$$

whose stationarity conditions solve in three steps:

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial \mathbf{T}_{jk}} &= \frac{n_{jk}}{\mathbf{T}_{jk}} - c_j = 0 && \text{(first-order condition)} \\ \mathbf{T}_{jk} &= \frac{n_{jk}}{c_j} && \text{(solve for } \mathbf{T}_{jk} \text{)} \\ c_j &= \sum_k n_{jk} = n_{j\cdot} && \text{(row constraint } \sum_k \mathbf{T}_{jk} = 1 \text{)} \\ \hat{\mathbf{T}}_{jk} &= \frac{n_{jk}}{n_{j\cdot}} && \text{(counted estimator)} \end{aligned} \quad (\text{S26})$$

The interior condition applies to cells with positive counts; for a cell with $n_{jk} = 0$ the log-likelihood is non-increasing in $\mathbf{T}_{jk}$, so its maximizer is the boundary value $\hat{\mathbf{T}}_{jk} = 0$, which the same formula returns. Multiplying $\hat{\mathbf{T}}_{jk}$ by the row frequency $n_{j\cdot}/(T-1)$ recovers $n_{jk}/(T-1)$, the empirical frequency of the state pair (j, k) . The counted estimator therefore reproduces the in-sample joint law of (S_{t-1}, S_t) exactly; everything it asserts about longer horizons is the Chapman–Kolmogorov extension $\mathbf{T}^\tau$ of that one-step law (Bilmes, 1998; Grinstead and Snell, 1997). The emission estimator completes the pair: within each state, μ_k and σ_k are the sample mean and standard deviation of the assigned observations, and σ_k enters Eq. (S18) as the scale, so the emission standard deviation is $\sigma_k \sqrt{\nu/(\nu-2)}$. The aggregate consequence of this scale convention is quantified in Eq. (S30).

The initial distribution ties the pieces together. Let n_k count the occupancy of state k over the T observations, and let $\hat{\pi}_k = n_k/T$ be the occupancy frequency. The row total satisfies $n_{j\cdot} = n_j - \mathbf{1}\{s_T = j\}$, and the column total satisfies $n_{\cdot k} = n_k - \mathbf{1}\{s_1 = k\}$, because only the last observation opens no transition and

only the first closes none. Substituting both identities gives:

$$\begin{aligned}
(\hat{\pi} \hat{\mathbf{T}})_k &= \sum_j \frac{n_j}{T} \cdot \frac{n_{jk}}{n_{j\cdot}} && \text{(definitions)} \\
&= \sum_j \frac{n_{j\cdot}}{T} \cdot \frac{n_{jk}}{n_{j\cdot}} + \sum_j \frac{\mathbf{1}\{s_T = j\}}{T} \cdot \frac{n_{jk}}{n_{j\cdot}} && (n_j = n_{j\cdot} + \mathbf{1}\{s_T = j\}) \\
&= \frac{n_{\cdot k}}{T} + \frac{\hat{\mathbf{T}}_{s_T, k}}{T} && \text{(sum transitions into } k\text{)} \\
&= \hat{\pi}_k + \frac{\hat{\mathbf{T}}_{s_T, k} - \mathbf{1}\{s_1 = k\}}{T}. && (n_{\cdot k} = n_k - \mathbf{1}\{s_1 = k\})
\end{aligned} \tag{S27}$$

The occupancy frequencies therefore satisfy the stationarity equation $\hat{\pi} \hat{\mathbf{T}} = \hat{\pi}$ up to a fixed-point residual bounded by $1/T$. For the fitted SPY model, the solved eigenvector $\bar{\pi}$ deviates from the occupancy frequencies by at most 3.7×10^{-4} , the same order as $1/T$, against occupancy values between 0.0058 and 0.017. The initial distribution is consequently not an abstract spectral object: $\bar{\pi}$ is, up to a correction of the order of $1/T$, the vector of in-sample state frequencies.

S4.3. Stationary distribution and moments

The stationary observation law of the no-jump generator follows directly from Eq. (S17). Marginalizing the state at a single time gives the mixture density and CDF:

$$p(g) = \sum_{k=1}^N \bar{\pi}_k f_k(g), \quad F_{\text{mix}}(x) = \sum_{k=1}^N \bar{\pi}_k \Psi_\nu \left(\frac{x - \mu_k}{\sigma_k} \right), \tag{S28}$$

where Ψ_ν is the CDF of the standard Student- t with ν degrees of freedom. Every ingredient of Eq. (S28) is an empirical summary: the weights are the in-sample state frequencies, and the component locations and scales are the within-state sample moments. The generator's marginal is therefore a 100-component semi-parametric estimate of the in-sample distribution, and reproducing the empirical marginal requires no further tuning. The mean is matched by construction. Weighting the within-state means by the occupancy frequencies telescopes into the grand mean:

$$\begin{aligned}
\sum_k \hat{\pi}_k \mu_k &= \sum_k \frac{n_k}{T} \cdot \frac{1}{n_k} \sum_{t: s_t = k} G_t && \text{(definitions of } \hat{\pi}_k, \mu_k\text{)} \\
&= \frac{1}{T} \sum_{t=1}^T G_t, && \text{(each } t \text{ counted once)}
\end{aligned} \tag{S29}$$

and replacing $\hat{\pi}$ by $\bar{\pi}$ changes the value only through the $1/T$ correction of Eq. (S27): the model mean is 0.0636 yr^{-1} against the sample mean of 0.0631 yr^{-1} . The distributional match is equally direct. Evaluated at the fitted SPY model,

the largest gap between the model CDF and the empirical CDF is $\sup_x |F_{\text{mix}}(x) - \hat{F}_T(x)| = 0.0045$, where $\hat{F}_T$ is the in-sample empirical CDF. The two-sample Kolmogorov–Smirnov (KS) critical value at significance level $\alpha = 0.05$ for two samples of 2,766 observations is $1.358\sqrt{2/T} = 0.0365$, so the model-versus-data distance consumes about one eighth of the rejection threshold. Rejections of simulated paths are therefore driven by path-level sampling variation rather than by a misplaced marginal, which is consistent with the in-sample pass rates in Supplementary Table S3.

The second moment makes the role of the emission scale convention explicit. Conditioning on the state and applying the law of total variance gives:

$$\begin{aligned} \text{Var}(G_t) &= \mathbb{E}[\text{Var}(G_t | S_t)] + \text{Var}(\mathbb{E}[G_t | S_t]) \quad (\text{law of total variance}) \\ &= \frac{\nu}{\nu-2} \sum_k \bar{\pi}_k \sigma_k^2 + \sum_k \bar{\pi}_k (\mu_k - \bar{\mu})^2, \quad (\text{Student-}t \text{ variance}) \end{aligned} \quad (\text{S30})$$

where $\bar{\mu} = \sum_k \bar{\pi}_k \mu_k$ is the stationary mean. Define the within-state variance $V_w = \sum_k \bar{\pi}_k \sigma_k^2$ and the between-state variance $V_b = \sum_k \bar{\pi}_k (\mu_k - \bar{\mu})^2$. The corresponding empirical decomposition of the sample variance is $V_b + V_w$, up to degrees-of-freedom corrections, so the model variance exceeds the sample variance by $2V_w/(\nu-2)$. At the SPY fit, $V_w = 0.17 \text{ yr}^{-2}$, only 3.6% of the total variance of 4.60 yr^{-2} : the quantile partition concentrates almost all dispersion between states rather than within them. The inflation is correspondingly small, a model variance of 4.72 against 4.60 yr^{-2} (1.3% in standard deviation, 2.17 against 2.14 yr^{-1}), and vanishes under Gaussian emissions (4.61 yr^{-2}).

The fourth moment explains the emission comparison in Supplementary Table S3. Write $G_t = \mu_k + \sigma_k X$ within state k , where X is a standard Student- t draw, and let $\delta_k = \mu_k - \bar{\mu}$. Expanding the fourth central moment and dropping the odd-power terms, which vanish by the symmetry of ψ_ν , gives:

$$\begin{aligned} \mathbb{E}[(G_t - \bar{\mu})^4] &= \sum_k \bar{\pi}_k \mathbb{E}[(\delta_k + \sigma_k X)^4] \quad (\text{condition on the state}) \\ &= \sum_k \bar{\pi}_k [\delta_k^4 + 6\delta_k^2 \sigma_k^2 \mathbb{E}(X^2) \\ &\quad + \sigma_k^4 \mathbb{E}(X^4)], \end{aligned} \quad (\text{S31})$$

with the standard Student- t moments, finite for $\nu > 4$:

$$\mathbb{E}(X^2) = \frac{\nu}{\nu-2}, \quad \mathbb{E}(X^4) = \frac{3\nu^2}{(\nu-2)(\nu-4)}. \quad (\text{S32})$$

At $\nu = 5$ these values are $5/3$ and 25 : the within-state kurtosis is 9 , against 3 for Gaussian emissions. The excess kurtosis of the stationary mixture is $\mathbb{E}[(G_t - \bar{\mu})^4] / \text{Var}(G_t)^2 - 3$; evaluating it at the SPY fit gives 8.34 with Student- t emissions and 5.82 with Gaussian emissions, against an observed value of 7.72 . The emission choice enters Eq. (S31) through both moments in Eq. (S32); the larger change by far is $\mathbb{E}(X^4)$ rising from 3 to 25 , which acts on the σ_k^4 terms

that the wide tail states dominate. The per-path simulated values reported in Supplementary Table S3 (7.5 and 5.5) sit below these population values because the sample kurtosis of a heavy-tailed series of 2,766 observations is biased downward, but the population gap of 2.5 between the emission choices matches the simulated gap of 2.0 in both direction and size. The mixture also carries the asymmetry of the fitted state occupancies without an explicit skewness parameter: simulating 1,000 in-sample-length paths gives a mean sample skewness of -0.76 for HMM-NJ and -0.73 for HMM-WJ, against the observed -0.75 (Supplementary Table S1). Because the states partition the Laplace quantiles of an asymmetric series, the occupancy-weighted mixture inherits that asymmetry through the μ_k and σ_k estimates.

S4.4. Dependence structure without jumps

The results stated so far concern single time slices; the dependence structure is where the no-jump generator falls short of the data, and the failure is structural rather than a matter of estimation. Within state k , the sojourn time is geometric: the chain remains for m steps with probability $\mathbf{T}_{kk}^{m-1}(1 - \mathbf{T}_{kk})$ and therefore stays for $1/(1 - \mathbf{T}_{kk})$ steps on average, between one and two days for every fitted state (Supplementary Fig. S3c). The autocorrelation of absolute returns inherits this short memory. Let $a_k = \mathbb{E}[|G_t| | S_t = k]$ be the within-state mean absolute return and $\bar{a} = \sum_k \bar{\pi}_k a_k$ its stationary average. Because observations are conditionally independent given the states, the lag- τ product moment reduces to a state-pair sum:

$$\begin{aligned} \mathbb{E}[|G_t| | G_{t+\tau}|] &= \sum_{j,k} \Pr(S_t = j, S_{t+\tau} = k) a_j a_k \quad (\text{conditional independence}) \\ &= \sum_{j,k} \bar{\pi}_j (\mathbf{T}^\tau)_{jk} a_j a_k, \quad (\text{Chapman-Kolmogorov}) \end{aligned} \tag{S33}$$

so the autocorrelation function (ACF) of $|G_t|$ is:

$$\text{ACF}_{|G|}(\tau) = \frac{\sum_{j,k} \bar{\pi}_j (\mathbf{T}^\tau)_{jk} a_j a_k - \bar{a}^2}{\text{Var}(|G_t|)}. \tag{S34}$$

As τ grows, $\mathbf{T}^\tau$ converges to its rank-one limit, every row equal to $\bar{\pi}$, at the geometric rate $|\theta_2|^\tau$, where θ_2 is the second-largest eigenvalue of $\mathbf{T}$ in modulus; the ACF is therefore bounded by a constant multiple of $|\theta_2|^\tau$. For the fitted SPY matrix, $|\theta_2| = 0.253$, and $|\theta_2|^\tau$ falls below 0.01 from lag 4 onward. Evaluating Eq. (S33) gives a model ACF of 0.26 at lag 1, close to the observed 0.30 because the counted matrix pins the one-step state-pair law, but 0.021 at lag 3 and effectively zero from lag 5, while the observed ACF remains near 0.26 at lag 5 and 0.14 at lag 25. No first-order chain fitted by transition counting can do otherwise: matching the one-step law fixes $\mathbf{T}$, and the spectral gap of that matrix then forces geometric mixing within a week. This is the quantitative content of the main-text statement that the counted transition probabilities

leave extreme-return states too quickly to reproduce volatility clustering, and it is the failure the jump mechanism corrects.

S4.5. *Jump-episode arithmetic*

The jump mechanism overlays the fast-mixing chain with rare, long episodes, and its operating characteristics follow from renewal arguments. On each step with no active episode, an episode with at least one forced step begins with probability $p_e = \epsilon(1 - e^{-\lambda})$, since a trigger whose Poisson draw returns zero is discarded; conditional on starting, the episode length is the zero-truncated Poisson draw with mean $\mathbb{E}[K \mid K \geq 1] = \lambda/(1 - e^{-\lambda})$. At $\lambda = 90$, the correction $e^{-\lambda}$ is below 10^{-39} , so $p_e = \epsilon$ and $\mathbb{E}[K \mid K \geq 1] = \lambda$ to machine precision. The long-run fraction of forced steps, denoted π_J , follows from the renewal-reward theorem. A renewal cycle consists of the free steps preceding an episode and the episode itself: the number of free trials up to and including the trigger is geometric with mean $1/p_e$, the triggering trial is itself the first forced step, and the episode contributes $\mathbb{E}[K \mid K \geq 1]$ forced steps in total:

$$\begin{aligned}
 \pi_J &= \frac{\mathbb{E}[K \mid K \geq 1]}{\frac{1 - p_e}{p_e} + \mathbb{E}[K \mid K \geq 1]} && \text{(forced steps over cycle length)} \\
 &= \frac{p_e \mathbb{E}[K \mid K \geq 1]}{1 - p_e + p_e \mathbb{E}[K \mid K \geq 1]} && \text{(multiply through by } p_e) \\
 &= \frac{\epsilon \lambda}{1 - \epsilon + \epsilon e^{-\lambda} + \epsilon \lambda} && (p_e \mathbb{E}[K \mid K \geq 1] = \epsilon \lambda) \\
 &\approx \frac{\epsilon \lambda}{1 + \epsilon \lambda} = 0.0089. && \text{(drop } \epsilon e^{-\lambda}, \text{ then } \epsilon) \tag{S35}
 \end{aligned}$$

Fewer than one step in a hundred is forced at the fitted values; the measured fraction across 400 simulated in-sample-length paths is 0.0088. The product $\epsilon \lambda$ is the quantity that sets this fraction, a point that matters for identification in Section S4.6.

At the path level, the same arithmetic determines how episodes are distributed across the ensemble. A simulated path contains no episode only if every one of its $T - 1$ transition steps fails to trigger one, so the jump-path probability at $T = 2,766$ is $1 - (1 - p_e)^{T-1} = 24.2\%$; the simulated ensemble gives 24.8% (99 of 400 paths), matching the $\sim 24\%$ jump-path fraction conditioned on in Supplementary Fig. S2. Episodes begin at the long-run rate of one per expected renewal cycle, so a path carries about 0.274 episodes in expectation. The expected gap between episodes is $1/\epsilon = 10^4$ trading days, about 40 trading years, so the ensemble is a two-population mixture: roughly three quarters of the paths never trigger an episode and follow the HMM-NJ law exactly, while the remaining quarter carry a single episode (about one jump path in eight carries a second). Within a jump path, the forced fraction is $\lambda/T = 3.3\%$ for a single episode; the simulated mean is 3.6%, slightly higher through the multi-episode paths. Forced steps differ sharply from free ones. Under the forced-step distribution u of Eq. (S19), the conditional standard deviation is 5.38 yr^{-1} against the

free-step 2.17 yr^{-1} , a variance ratio of 6.1, and the mean absolute return is $\bar{a}_J = \sum_k u_k a_k = 4.96 \text{ yr}^{-1}$ against $\bar{a} = 1.47 \text{ yr}^{-1}$. The stationary state marginal of the augmented chain is the mixture $(1 - \pi_J)$ times the free-phase distribution plus π_J times u ; the free phase relaxes from the post-episode tail state back to $\bar{\pi}$ at the rate $|\theta_2|$ within a few steps. Using $\bar{\pi}$ as the initial law therefore misallocates $O(\pi_J) \approx 1\%$ of probability mass, which quantifies the main-text remark that $\bar{\pi}$ is exact for the ordinary transition matrix but approximate for the jump-augmented process.

S4.6. Autocorrelation and kurtosis under jumps

The tuning objective of Eq. (3) scores the absolute-return ACF and the excess kurtosis, and both statistics admit closed-form approximations that show what (ϵ, λ) control. The ACF contribution rests on one more renewal quantity: the probability $h(\tau)$ that, given a step is forced, the step τ later belongs to the same episode. A uniformly chosen forced step lands in an episode with the length-biased law, an episode of length m covering m forced steps, and its position within that episode is uniform. Combining the two gives:

$$\begin{aligned} h(\tau) &= \sum_{m \geq 1} \frac{m \Pr(K = m \mid K \geq 1)}{\mathbb{E}[K \mid K \geq 1]} \cdot \frac{(m - \tau)_+}{m} \\ &= \frac{\mathbb{E}[(K - \tau)_+ \mid K \geq 1]}{\mathbb{E}[K \mid K \geq 1]} && \text{(sum over } m) \\ &\approx \left(1 - \frac{\tau}{\lambda}\right)_+, && (K \text{ concentrated near } \lambda) \end{aligned} \tag{S36}$$

where $(x)_+ = \max(x, 0)$ and the approximation holds because the Poisson length concentrates within a few multiples of $\sqrt{\lambda} \approx 9.5$ around λ . The same-episode probability therefore decays almost linearly from one to zero over the horizon λ . The ACF follows by conditioning on the forced-step indicator $Z_t = \mathbf{1}\{\text{step } t \text{ is forced}\}$. Given the indicator path, every forced step's state is an independent draw from u , so two forced observations are conditionally independent; the dependence linking a forced observation to a later free one runs only through the episode's final state and, like the free-phase covariance itself, decays at the $|\theta_2|^\tau$ rate of Section S4.4. For lags beyond that scale, only the conditional means covary. Write $A_t = |G_t|$ for the absolute return. The conditional mean is $\mathbb{E}[A_t \mid Z_t] = \bar{a} + Z_t (\bar{a}_J - \bar{a})$, and for rare episodes the joint forced probability is $\Pr(Z_t = 1, Z_{t+\tau} = 1) \approx \pi_J h(\tau)$, so:

$$\begin{aligned} \text{Cov}(A_t, A_{t+\tau}) &\approx \text{Cov}(\mathbb{E}[A_t \mid Z_t], \mathbb{E}[A_{t+\tau} \mid Z_{t+\tau}]) && \text{(independence given } Z) \\ &= (\bar{a}_J - \bar{a})^2 \text{Cov}(Z_t, Z_{t+\tau}) && \text{(conditional mean in } Z) \\ &\approx (\bar{a}_J - \bar{a})^2 \pi_J h(\tau), && \text{(episodes rare)} \end{aligned} \tag{S37}$$

so the jump mechanism adds an ACF component of height $\pi_J(\bar{a}_J - \bar{a})^2 / \text{Var}(A_t)$ that decays almost linearly to zero at lag λ , precisely the slow component the

Table S2: **Episode-overlap prediction of the jump-path absolute-return autocorrelation function (ACF) against simulation.** The prediction evaluates Eq. (S37) at the fitted SPY values with the single-episode forced fraction λ/T in place of π_J and the same-episode probability $h(\tau)$ of Eq. (S36). The simulated column is the mean ACF of $|G_t|$ over the 99 jump-active paths among 400 simulated in-sample-length paths. The free-phase component of Eq. (S33) is excluded from the prediction and has decayed below 0.01 at the lags shown.

Lag τ	$h(\tau)$	Predicted ACF	Simulated ACF
5	0.944	0.124	0.137
10	0.889	0.117	0.123
15	0.833	0.110	0.115
20	0.778	0.102	0.105
25	0.722	0.095	0.099

counted chain cannot produce. For a jump path carrying a single episode, the relevant forced fraction is λ/T rather than π_J , and evaluating the resulting expression at the fitted values, with $\text{Var}(A_t) = 3.01$ under the same conditioning, gives a predicted plateau of 0.13. Supplementary Table S2 compares this prediction with the mean simulated jump-path ACF: the two differ by 0.013 at lag 5 and by less than 0.01 from lag 10 onward. At lags 1 to 3 the simulated values exceed the episode component (0.35 against 0.13 at lag 1) because the free-phase geometric component of Section S4.4 adds on top before it dies out.

The kurtosis contribution pulls in the opposite direction. Treating a jump path as a two-regime mixture with weights $(1 - \lambda/T, \lambda/T)$ and evaluating the fourth-moment expression of Eq. (S31) with the forced-step distribution u in place of $\bar{\pi}$ for the episode regime gives a population excess kurtosis of 7.20 for a single-episode jump path, against 8.34 with no jumps. Forced steps inflate the variance, by about 17% at the fitted values, faster than they inflate the fourth moment, because u draws uniformly within each tail set rather than extending the extreme tail. Adding jump content therefore lowers kurtosis toward and past the observed value while it raises the slow ACF component; the jump-share sweep reported in the main text traces the same tradeoff empirically, with the Hill tail index flat across the sweep because the tail exponent is set by ν , not by the jump content. The objective of Eq. (3) balances these two effects.

The same expressions show how the two jump parameters are identified, and where they are not. Within a single-episode jump path, the ACF component of Eq. (S37) has height proportional to λ/T and horizon λ , and the kurtosis shift depends on the same forced fraction: every statistic entering Eq. (3) for such a path is controlled by λ . The parameter ϵ moves the objective only through the incidence of jump paths, which the jump-active conditioning removes, and through the small fraction of multi-episode paths. This weak dependence is consistent with the tuned $\epsilon = 10^{-4}$ lying at the lower edge of the explored range for SPY and for all three cross-asset tickers, and with the Monte Carlo sensitivity of the tuned optima noted in Section S8. The objects entering every expression in this appendix, the fitted boundaries, transition matrix, and residence times, are plotted in Supplementary Fig. S3.




Figure S3: **Fitted hidden Markov model with jumps (HMM-WJ) internals for SPY with $N = 100$ states.** Panel (a) overlays the fitted Laplace cumulative distribution function (CDF) on the empirical CDF; the 99 vertical lines mark the equal-probability quantile boundaries Q_k that define the hidden states. Panel (b) displays the estimated transition matrix $\mathbf{T}$ in $\log_{10}$ color scale; the dominant near-diagonal band reflects short-range regime persistence. Panel (c) plots the expected natural residence time $1/(1 - T_{kk})$ for each state on a $\log_{10}$ scale, with the bottom tail set $\mathcal{S}_-$ (states 1–5, red) and the top tail set $\mathcal{S}_+$ (states 96–100, teal) highlighted. The dashed horizontal line marks the fitted mean jump duration $\lambda = 90$ steps, quantifying the two-order-of-magnitude gap between natural residence times and the jump-duration override.

S5. Transition matrix internals and partition diagnostics

To determine why the ordinary state process required a duration override, we inspected the fitted hidden Markov model with jumps (HMM-WJ) components for SPY at $N = 100$ (Supplementary Fig. S3). The fitted Laplace CDF and the empirical CDF are shown with the 99 state boundaries Q_k . The transition matrix $\mathbf{T}$ has a near-diagonal band, with most off-diagonal mass within ± 5 states of the diagonal. Under the ordinary Markov transitions, every state has an expected residence time of only 1–2 steps. The fitted jump duration is $\lambda = 90$ steps, about two orders of magnitude longer. Every state had at least one in-sample observation. This residence-time gap, rather than missing states, motivated the duration override.

S6. Emission distribution and semi-Markov comparison

To compare the emission and duration specifications, we evaluated three variants on SPY (Supplementary Table S3): a hidden semi-Markov model (HSMM) baseline adapted from Bulla and Bulla (2006) ($K = 8$ states, negative-binomial dwell times with a geometric fallback, and Student- t emissions), HMM-WJ with the original Gaussian emissions ($N = 100$), and HMM-WJ with Student- t emissions ($N = 100$). Both HMM-WJ variants had higher in-sample KS and AD pass rates than the HSMM-style baseline. The HMM-WJ variants use 100 states and the HSMM uses 8, so this comparison does not assign the pass-rate difference solely to the duration model. Holding the rest of HMM-WJ fixed, changing the emission from Normal to Student- t ($\nu = 5$) increased simulated excess kurtosis from 5.5 to 7.5, against an observed value of 7.7. The

Table S3: **Three-way comparison of emission distributions and duration specifications for SPY**. Standard errors in parentheses; bold marks the best value per row. The HSMM is a hidden semi-Markov-style model adapted from Bulla and Bulla (2006) ($K = 8$ states, negative-binomial dwell times with a geometric fallback, and Student- t emissions); HMM-WJ: hidden Markov model with jumps ($N = 100$). All entries are computed from 1,000 simulated paths at significance level $\alpha = 0.05$. Kolmogorov–Smirnov (KS), Anderson–Darling (AD), and mean absolute error in the autocorrelation function (ACF-MAE) are defined in Section S3.

Metric	HSMM (Student- t)	HMM-WJ (Gaussian)	HMM-WJ (Student- t)
States	$K = 8$	$N = 100$	$N = 100$
<i>In-sample (2,766 trading days, 2014–2024)</i>			
KS pass rate (%)	82.0 (1.2)	97.0 (0.5)	98.3 (0.4)
AD pass rate (%)	42.5 (1.6)	90.5 (0.9)	94.8 (0.7)
Excess kurtosis	4.8 (0.08)	5.5 (0.03)	7.5 (0.09)
ACF-MAE	0.059 (<0.001)	0.052 (<0.001)	0.053 (<0.001)
Wasserstein-1	0.176 (0.001)	0.101 (0.002)	0.097 (0.001)
Hellinger	0.113 (<0.001)	0.076 (<0.001)	0.074 (<0.001)
<i>Out-of-sample (249 trading days, 2025)</i>			
KS pass rate (%)	96.2 (0.6)	95.0 (0.7)	95.4 (0.7)
AD pass rate (%)	96.7 (0.6)	95.9 (0.6)	95.3 (0.7)
Excess kurtosis	4.1 (0.13)	4.9 (0.08)	6.2 (0.15)
ACF-MAE	0.042 (<0.001)	0.040 (<0.001)	0.040 (<0.001)
Wasserstein-1	0.287 (0.002)	0.275 (0.005)	0.275 (0.005)
Hellinger	0.239 (0.001)	0.212 (0.001)	0.210 (0.001)

Student- t emission also increased the KS and AD pass rates, while ACF-MAE and the distributional distances changed modestly. The emission choice therefore explained the kurtosis gain in this controlled HMM-WJ comparison, while it changed temporal fit little.

S7. State-resolution sensitivity

To test whether the results depended on state resolution, we swept the partition over $N \in \{30, 60, 90, 100, 150, 200, 350\}$ and re-evaluated the in-sample Kolmogorov–Smirnov (KS) and Anderson–Darling (AD) pass rates and the mean absolute error in the absolute-return autocorrelation function (ACF-MAE) for the hidden Markov model with no jumps (HMM-NJ) and the hidden Markov model with jumps (HMM-WJ), holding the tail-set size N_{tail} fixed (Supplementary Table S6). HMM-NJ was insensitive to the partition: its pass rates moved by less than one percentage point and its ACF-MAE was unchanged at the reported precision, as expected for a model with no duration override whose behavior could depend on the width of a tail state. HMM-WJ traded the two metrics against each other. Because N_{tail} is fixed, a finer partition makes the tail set a narrower quantile band, so each forced step lands further out, and ACF-MAE

Table S4: **Hybrid composition method broken out by branch.** N_{tickers} : number of tickers assigned to each branch by the $R^2_{\text{preserve}} = 0.80$ threshold. The clipped variance-preserving branch does not trigger anywhere in the universe ($\rho_i < 0.90$ for all tickers) and is omitted. The two trackers in the R^2 -preserving branch (QQQ and SPYG) had median absolute errors of 0.005 for β and 0.001 for R^2 , with a KS pass rate of 100%.

Branch	N_{tickers}	$ \hat{\beta} - \beta $	$ \hat{R}^2 - R^2_{\text{real}} $	KS pass (%)	W_1	κ
hybrid	421	0.018	0.021	73.5	0.250	7.142
r2-preserve	2	0.005	0.001	100.0	0.099	11.525

Table S5: **Synthetic-tracker grid for the R^2 -preserving branch.** The observed universe places only two tickers on this branch, so we generated synthetic trackers at each $(\beta_{\text{true}}, R^2_{\text{true}})$ cell, composed them with Algorithm 2, and recovered $(\hat{\beta}, \hat{R}^2)$ by ordinary least squares. Entries are medians over 100 replications per cell with cross-replication standard deviations in parentheses; the Kolmogorov–Smirnov (KS) pass rate is the fraction of replications whose two-sample test against the target marginal exceeds $p = 0.05$. Every cell was assigned to the R^2 -preserving branch.

β_{true}	R^2_{true}	$\hat{\beta}$	$\hat{R}^2$	KS pass (%)
0.8	0.80	0.8006 (0.0072)	0.7998 (0.0044)	100.0
0.8	0.85	0.8006 (0.0059)	0.8497 (0.0041)	100.0
0.8	0.90	0.8004 (0.0048)	0.8999 (0.0028)	100.0
0.8	0.95	0.7995 (0.0033)	0.9500 (0.0013)	100.0
0.8	0.99	0.7999 (0.0016)	0.9901 (0.0003)	100.0
1.0	0.80	1.0003 (0.0089)	0.8017 (0.0052)	100.0
1.0	0.85	0.9994 (0.0084)	0.8497 (0.0039)	100.0
1.0	0.90	0.9999 (0.0065)	0.8999 (0.0026)	100.0
1.0	0.95	1.0005 (0.0036)	0.9502 (0.0013)	100.0
1.0	0.99	1.0000 (0.0018)	0.9900 (0.0003)	100.0
1.2	0.80	1.2007 (0.0126)	0.8001 (0.0048)	100.0
1.2	0.85	1.2011 (0.0105)	0.8503 (0.0043)	100.0
1.2	0.90	1.1996 (0.0081)	0.8994 (0.0027)	100.0
1.2	0.95	1.1995 (0.0052)	0.9498 (0.0014)	100.0
1.2	0.99	1.1998 (0.0025)	0.9900 (0.0003)	100.0

Table S6: **Sensitivity of in-sample fit to the number of states N .** All entries are computed from 1,000 simulated in-sample-length paths for SPY at significance level $\alpha = 0.05$, with the tail-set size N_{tail} and the remaining hyperparameters of Table 1 held fixed and the (ϵ, λ) grid search re-run at each N for the hidden Markov model with jumps (HMM-WJ). Standard errors in parentheses: binomial for pass rates, bootstrap ($B = 500$) for the mean absolute error in the absolute-return autocorrelation function (ACF-MAE). HMM-NJ denotes the hidden Markov model with no jumps. Values at $N = 100$ reproduce the corresponding entries of Table 2.

Model	N	KS pass (%)	AD pass (%)	ACF-MAE
HMM-NJ	30	99.3 (0.3)	97.8 (0.5)	0.059 (<0.001)
HMM-NJ	60	99.5 (0.2)	98.5 (0.4)	0.059 (<0.001)
HMM-NJ	90	99.3 (0.3)	98.6 (0.4)	0.059 (<0.001)
HMM-NJ	100	99.3 (0.3)	98.2 (0.4)	0.059 (<0.001)
HMM-NJ	150	99.1 (0.3)	98.7 (0.4)	0.059 (<0.001)
HMM-NJ	200	99.2 (0.3)	98.3 (0.4)	0.059 (<0.001)
HMM-WJ	30	98.4 (0.4)	97.1 (0.5)	0.057 (<0.001)
HMM-WJ	60	96.7 (0.6)	91.4 (0.9)	0.053 (<0.001)
HMM-WJ	90	98.3 (0.4)	95.2 (0.7)	0.055 (<0.001)
HMM-WJ	100	98.3 (0.4)	94.8 (0.7)	0.053 (<0.001)
HMM-WJ	150	97.7 (0.5)	91.7 (0.9)	0.052 (<0.001)
HMM-WJ	200	96.0 (0.6)	82.8 (1.2)	0.045 (0.001)

improved from 0.057 at $N = 30$ to 0.045 at $N = 200$. The same narrowing left fewer observations behind each tail state’s emission parameters, and the AD pass rate, which weights tail discrepancies heavily, moved the other way, from 97.1% at $N = 30$ to 82.8% at $N = 200$. The intermediate values were not monotone, which is consistent with re-running the (ϵ, λ) grid search on a finite path ensemble at each N . At $N = 350$ several quantile bins were never visited in the historical record, leaving the corresponding emission parameters undefined. We adopted $N = 100$ for all subsequent experiments as a midpoint that held the KS pass rate above 98% and the AD pass rate within a few percentage points of its best value while capturing part of the ACF-MAE improvement.

S8. Cross-asset generalization

To test whether the single-asset procedure generalized beyond SPY, we repeated the fitting and evaluation on three stocks: NVDA (semiconductors, high volatility), JNJ (health care, low volatility), and JPM (financials, moderate volatility). Per-ticker tune outputs at ($n_{\text{paths}} = 200, \text{seed} = 1234$) were reproducible with these settings (Supplementary Table S9). The objective combines ACF and kurtosis errors with fixed weights. Because kurtosis varies across finite samples of heavy-tailed paths, the selected grid point can change with the Monte Carlo sample. The reported (ϵ^*, λ^*) values are reproducible with the stated path count and seed, but they are not unique per-ticker optima. At the selected parameter values, in-sample HMM-WJ KS pass rates were 98.9% for NVDA, 99.4% for JNJ, and 99.1% for JPM. Out-of-sample KS pass rates on the 249-day




Figure S4: **Branch assignment in calibrated $(\beta, R^2_{\text{real}})$ space.** The $R^2_{\text{preserve}} = 0.80$ threshold (dashed line) separates tracker assets from single stocks. In this universe, only QQQ and SPYG exceed the threshold, and both are assigned to the R^2 -preserving branch. The remaining 421 tickers populate the variance-preserving branch; the clipped variance-preserving branch is unused because the market loading ratio $\rho_i = \beta_i^2 \sigma_m^2 / \sigma_{\text{gen},i}^2$ stays below 0.90 for every ticker.

Table S7: **Aggregate metrics by β quartile, across composition methods.** Quartile cutoffs are computed from the calibrated β distribution across the 423-ticker universe. The hybrid KS pass rate declines from 87.6% in Q1 to 65.9% in Q4 as the market loading ratio ρ_i grows; the naive KS pass rate falls from 24.6% to 1.3% over the same range.

β bucket	Method	$ \hat{\beta} - \beta $	KS pass (%)	W_1	κ	$\text{Var}(g)$
Q1 (low β)	Naive	0.018	24.6	0.403	8.971	18.1
	Gaussian SIM	0.015	0.0	0.620	0.622	14.0
	Hybrid	0.016	87.6	0.177	8.369	14.6
	JumpHMM-on-residuals	0.015	73.2	0.204	7.059	14.3
	Block bootstrap	0.015	71.3	0.208	6.286	13.9
	GARCH(1,1)- t	0.015	57.8	0.253	5.255	13.4
Q2	Naive	0.021	1.9	0.621	7.510	25.1
	Gaussian SIM	0.016	0.2	0.653	1.509	18.0
	Hybrid	0.017	74.9	0.231	6.889	19.0
	JumpHMM-on-residuals	0.017	82.7	0.200	7.799	18.7
	Block bootstrap	0.016	79.8	0.197	7.213	18.0
	GARCH(1,1)- t	0.015	76.4	0.227	6.470	17.8
Q3	Naive	0.023	0.7	0.792	7.390	33.3
	Gaussian SIM	0.017	1.4	0.688	1.947	23.4
	Hybrid	0.018	66.1	0.279	6.904	24.4
	JumpHMM-on-residuals	0.018	78.3	0.242	7.454	24.0
	Block bootstrap	0.017	75.0	0.236	6.888	22.8
	GARCH(1,1)- t	0.016	71.0	0.263	6.265	21.9
Q4 (high β)	Naive	0.029	1.3	0.965	7.233	51.2
	Gaussian SIM	0.022	0.6	0.825	2.054	36.3
	Hybrid	0.022	65.9	0.339	6.739	37.1
	JumpHMM-on-residuals	0.022	69.1	0.313	6.842	37.0
	Block bootstrap	0.021	67.2	0.311	6.334	35.3
	GARCH(1,1)- t	0.021	64.5	0.361	6.064	34.7

Table S8: **Clipping-branch stress test under market-volatility scaling.** The market growth-rate path was rescaled by $\gamma \in \{1, 2, 3\}$ with the calibrated marginals held fixed, and the full hybrid composition method re-evaluated on the 423-ticker universe. As γ rises, the share of clipped tickers grows.

γ	Clipped (%)
1	0.0
2	76.4
3	95.2

Table S9: **Cross-asset generalization of the hidden Markov model with jumps (HMM-WJ) across three single-stock tickers spanning sector and volatility regimes.** Per-ticker tune outputs (ϵ^*, λ^*) at $n_{\text{paths}} = 200$ and seed = 1234, and the corresponding HMM-WJ Kolmogorov–Smirnov (KS) and Anderson–Darling (AD) pass rates on 1,000 simulated paths at $\alpha = 0.05$ against each ticker’s in-sample (IS; 2014–2024) and out-of-sample (OoS; 2025) growth-rate history. Values in parentheses are binomial standard errors (SEs) for pass rates. The (ϵ^*, λ^*) values are reported as the parameter values used for the IS / OoS metrics rather than as canonical per-ticker optima.

Ticker	ϵ^*	λ^*	IS KS (%)	IS AD (%)	OoS KS (%)	OoS AD (%)
NVDA (semiconductors)	10^{-4}	70	98.9 (0.3)	96.7 (0.6)	97.7 (0.5)	96.1 (0.6)
JNJ (health care)	10^{-4}	70	99.4 (0.2)	96.3 (0.6)	74.0 (1.4)	61.3 (1.5)
JPM (financials)	10^{-4}	80	99.1 (0.3)	95.7 (0.6)	79.2 (1.3)	77.7 (1.3)

2025 hold-out were 97.7%, 74.0%, and 79.2%, respectively. Thus the strong in-sample fit transferred unevenly across assets, especially on the short holdout.

S9. SIM composition derivations and properties verification

To verify the two closed-form targets used by the variance-corrected single-index composition, we derived the variance scale of Eq. (8), the residual-variance target of Eq. (11), and the resulting large- T OLS coefficients and composed variance. To derive Eq. (8), first center the full-return draw as in Eq. (6), then set $\varepsilon_i = s_i \tilde{g}_i^c$ for $s_i \in [0, 1]$ and require that the composed series have the same variance as the generator path:

$$\text{Var}(g_i) = \text{Var}(\alpha_i + \beta_i g_m + \varepsilon_i) = \text{Var}(\beta_i g_m + \varepsilon_i) = \sigma_{\text{gen},i}^2, \quad (\text{S38})$$

where α_i drops out because it is constant. Expanding the sum gives:

$$\text{Var}(\beta_i g_m + \varepsilon_i) = \beta_i^2 \text{Var}(g_m) + 2\beta_i \text{Cov}(g_m, \varepsilon_i) + \text{Var}(\varepsilon_i). \quad (\text{S39})$$

Substituting $\varepsilon_i = s_i \tilde{g}_i^c$ gives $\text{Var}(\varepsilon_i) = s_i^2 \sigma_{\text{gen},i}^2$ and $\text{Cov}(g_m, \varepsilon_i) = s_i \text{Cov}(g_m, \tilde{g}_i^c)$. Under the assumption $\text{Cov}(\tilde{g}_i^c, g_m) \approx 0$, the cross term vanishes and the variance condition becomes $\beta_i^2 \sigma_m^2 + s_i^2 \sigma_{\text{gen},i}^2 = \sigma_{\text{gen},i}^2$. Solving for s_i^2 gives Eq. (8).

To derive Eq. (11), use the same zero-cross-covariance assumption and write the population R^2 identity for Eq. (4):

$$R^2 = \frac{\beta_i^2 \sigma_m^2}{\beta_i^2 \sigma_m^2 + \sigma_{\varepsilon,i}^2}. \quad (\text{S40})$$




Figure S5: **Calibrated R^2 distribution across the 423 non-SPY tickers.** Histogram of ordinary least-squares (OLS) R^2 against SPY over the 2014–2024 window. The two tickers crossing the $R^2_{\text{preserve}} = 0.80$ threshold (dashed line) are index exchange-traded funds (QQQ, SPYG); the highest- R^2 individual stock (BLK, 0.645) falls ~ 0.16 below the threshold. The R^2 -preserving branch therefore contains only QQQ and SPYG in this universe.




Figure S6: **Heavy-tail preservation: excess kurtosis and Hill tail index vs calibrated β .** Each point is a per-ticker median over 100 replications. *Left:* excess kurtosis, y -axis clipped to $[-2, 20]$ for readability. The dashed line marks the median real-data excess kurtosis across the universe ($\kappa_{\text{real}} \approx 13$); the hybrid and naive composition methods (orange and blue) track the generator's kurtosis, while the Gaussian single-index model (SIM) collapses to $\kappa \approx 1$. *Right:* Hill tail-index estimator on the upper 5% of $|g|$. The hybrid and naive medians remain near 0.37 throughout the β range; Gaussian SIM drops to ~ 0.22 . The hybrid's median being below κ_{real} reflects the JumpHMM marginal's own kurtosis rather than a composition-rule distortion: hybrid tracks naive (both retain the generator's tails) while Gaussian SIM does not.

Solving this identity for the residual variance gives Eq. (11). Because $\bar{\varepsilon}_i = 0$ after centering and scaling, the sample mean of a composed path obeys $\bar{g}_i = \alpha_i + \beta_i^{\text{eff}} \bar{g}_m$. Hence, the finite-sample OLS intercept was:

$$\hat{\alpha}_i = \bar{g}_i - \hat{\beta}_i \bar{g}_m = \alpha_i - (\hat{\beta}_i - \beta_i^{\text{eff}}) \bar{g}_m. \quad (\text{S41})$$

Under the zero-cross-covariance assumption, in the large- T limit OLS therefore gives $\hat{\alpha}_i \rightarrow \alpha_i$ and $\hat{\beta}_i \rightarrow \beta_i^{\text{eff}}$. Here $\beta_i^{\text{eff}} = \beta_i$ unless Eq. (10) clips the market loading. On the variance-preserving branch, the composed variance is $\sigma_{\text{gen},i}^2$. Without clipping, $\text{Var}(g_i) = \beta_i^2 \sigma_m^2 + s_i^2 \sigma_{\text{gen},i}^2 = \sigma_{\text{gen},i}^2$. Under clipping, Eq. (10) sets β_i^{eff} to give the same variance. On the R^2 -preserving branch, the composed variance is $\beta_i^2 \sigma_m^2 / R_{i,\text{real}}^2$, the value implied by the calibrated $(\beta_i, R_{i,\text{real}}^2)$ rather than by the generator marginal. On the ordinary-stock branch, the floor guarantees only that $\text{Var}(\varepsilon_i) \geq f \sigma_{\text{gen},i}^2$; it does not guarantee that adding the market term preserves the generator's tail distribution. The KS pass rate fell from 87.6% in the lowest-beta quartile to 65.9% in the highest-beta quartile (Supplementary Table S7).

S9.1. Role of the centered full-return draw as the residual

The hybrid composition places $\varepsilon_i = s_i \tilde{g}_i^c$, a centered and rescaled draw of asset i 's full return, in the residual slot of Eq. (4). On its face this is a category mismatch: the residual of a single-index decomposition is the idiosyncratic component of the return, orthogonal to the market, while $\tilde{g}_i^c$ is distributed like the centered total return, market-driven component included. This appendix states what the substitution preserves exactly, what it preserves approximately, and how the approximation degrades as the market loading grows.

The residual slot must supply three things: a zero mean, so that the intercept and market term alone set the conditional mean; the variance $\sigma_{\text{gen},i}^2 - \beta_i^2 \sigma_m^2$ not already contributed by the market term; and the asset-specific distributional shape, the heavy tails and volatility clustering that motivated the single-asset generator. Centering meets the first requirement exactly (Eq. (6)), the scale of Eq. (8) meets the second exactly, and the large- T limits $\hat{\alpha}_i \rightarrow \alpha_i$ and $\hat{\beta}_i \rightarrow \beta_i^{\text{eff}}$ hold because the draw is generated independently of the simulated market path (Sec. S1), so the recovered factor structure is exact as well. Any zero-mean, correctly scaled, market-independent draw satisfies the first two requirements; the full-return draw is chosen for the third, because it carries the per-asset marginal that the single-asset validation certified. Only the shape requirement is met approximately, and two exact statements bracket the approximation.

At $\rho_i = 0$ the substitution is not an approximation: Eq. (8) gives $s_i = 1$, the market term contributes no variance, and the composed series is the generator draw itself, so the composed distribution equals the validated single-asset marginal. For $\rho_i > 0$ the composed return is the sum of two independent centered draws, and the survival of the generator's shape follows in closed form. In population, conditional on the fitted parameters, write the deviation of the composed return from its mean as $X_m + X_i$, where $X_m = \beta_i g_m^c$ is the market contribution, g_m^c is the market path centered at its mean, and $X_i = s_i \tilde{g}_i^c$ is

the scaled asset draw; their variances are $v_m = \rho_i \sigma_{\text{gen},i}^2$ and $v_i = (1 - \rho_i) \sigma_{\text{gen},i}^2$. Independence factors every cross term in the fourth moment:

$$\begin{aligned}
\mathbb{E}[(X_m + X_i)^4] &= \mathbb{E}(X_m^4) + 4\mathbb{E}(X_m^3)\mathbb{E}(X_i) \\
&\quad + 6\mathbb{E}(X_m^2)\mathbb{E}(X_i^2) + 4\mathbb{E}(X_m)\mathbb{E}(X_i^3) \\
&\quad + \mathbb{E}(X_i^4) \quad (\text{independence}) \\
&= \mathbb{E}(X_m^4) + 6v_m v_i + \mathbb{E}(X_i^4). \quad (\mathbb{E}X_m = \mathbb{E}X_i = 0)
\end{aligned} \tag{S42}$$

Writing κ_m and κ_{gen} for the excess kurtosis of the market path and of the generator draw, the component fourth moments are $\mathbb{E}(X_m^4) = (\kappa_m + 3)v_m^2$ and $\mathbb{E}(X_i^4) = (\kappa_{\text{gen}} + 3)v_i^2$, where rescaling by s_i leaves excess kurtosis unchanged. The excess kurtosis of the composed return follows by dividing Eq. (S42) by the squared total variance:

$$\begin{aligned}
\kappa(X_m + X_i) &= \frac{(\kappa_m + 3)v_m^2 + 6v_m v_i + (\kappa_{\text{gen}} + 3)v_i^2}{(v_m + v_i)^2} - 3 \\
&= \frac{\kappa_m v_m^2 + \kappa_{\text{gen}} v_i^2 + 3(v_m + v_i)^2}{(v_m + v_i)^2} - 3 \quad (\text{collect terms}) \\
&= \kappa_m \rho_i^2 + \kappa_{\text{gen}} (1 - \rho_i)^2. \quad (v_m + v_i = \sigma_{\text{gen},i}^2)
\end{aligned} \tag{S43}$$

The generator's excess kurtosis survives composition with weight $(1 - \rho_i)^2$ and the market path's enters with weight ρ_i^2 ; a simulation check with independent generator draws matched Eq. (S43) to within 0.03 in mean excess kurtosis at $\rho_i \in \{0.1, 0.3, 0.5\}$. For a low-loading asset the composed tails are essentially the generator's: at $\rho_i = 0.10$ the generator's kurtosis retains weight 0.81. As ρ_i grows the composed shape moves away from the asset's own marginal, which is the pattern in the documented diagnostics: the hybrid tracks the generator's kurtosis across the β range while the Gaussian SIM does not (Supplementary Fig. S6), and the hybrid KS pass rate declines from the lowest to the highest β quartile (Supplementary Table S7).

The residual-fit methods provide the exact-orthogonal alternative: they fit a generator to each asset's OLS residual series, contain no market component to double count, and are compared with the hybrid in Table 3; fitting JumpHMM to residuals requires converting each residual series to a synthetic price path by cumulative compounding and a second fit per asset. In summary, $\tilde{g}_i^c$ is not a residual in the orthogonal-decomposition sense; it is a carrier of the validated marginal shape. The substitution fixes the mean, the variance, and the recovered factor loadings exactly, reproduces the asset's marginal exactly in the $\rho_i \rightarrow 0$ limit, and dilutes the generator's shape with weight $(1 - \rho_i)^2$ as the market share grows; the clipped and R^2 -preserving branches of Algorithm 2 handle the region where the dilution is largest.

S10. Uncertainty in the multi-asset jump comparison

The jump comparison used four seeds and 1,000 market replications per window, with all fits frozen (Table 6). Within each setting, naive and corrected composition shared the same asset and market draws. The no-jump and market-only settings reused the no-jump asset draws; the market-only and market-plus-asset settings reused the jump-enabled market draws. Enabled and disabled models used separate random-number streams, so matching replication identifiers did not imply identical innovations. Primary summaries included every path; jump-active subsets only described the mechanism. For each metric, we averaged across tickers within each replication and formed the difference between settings (Table S10). The Monte Carlo standard error was the sample standard deviation of the 1,000 differences divided by $\sqrt{1,000}$, and intervals used the mean difference plus or minus 1.96 standard errors (Table S10). Keeping the cross-section together retained dependence from the shared market draw. The intervals did not treat assets as independent market histories or include uncertainty from fitting or historical-window selection. Kolmogorov–Smirnov (KS) non-rejection remained descriptive because returns were temporally dependent. The temporal contrast between training and holdout remained clear in the replication-level uncertainty estimates (Table S10). The episode frequencies matched the rare-event setting: 247 of 1,000 training-length SPY paths and 22 of 1,000 holdout-length paths contained an episode. The corresponding shares across enabled asset generators were 24.24% and 2.46%. Market-only jumps reduced average 25-lag autocorrelation error for 99.5% of training assets but only 2.6% of holdout assets. Enabling jumps in both components reduced that error for 96.9% and 0.5%, respectively. These fractions described ticker-level mean changes, not significance tests. The temporal contrast therefore extended across most assets, within the limits of the short holdout and transferred SPY calibration.

Table S10: **Monte Carlo uncertainty in the effect of jump settings under corrected composition.** Differences are relative to jumps disabled, with 1,000 market replications per window. ΔACF_{25} is the change in mean absolute error of the absolute-return autocorrelation function through lag 25; negative values indicate improved temporal fit. ΔKS is the change in Kolmogorov–Smirnov non-rejection rate in percentage points. Brackets contain approximate 95% Monte Carlo (MC) intervals conditional on the fitted models and observed histories.

Jump setting	ΔACF_{25} [95% MC interval]	ΔKS [95% MC interval]
<i>Training: 2014–2024, 423 assets, 2,766 days</i>		
Market only	−0.00725 [−0.00812, −0.00638]	−0.30 [−1.08, 0.49]
Market + assets	−0.01988 [−0.02063, −0.01913]	−12.25 [−13.02, −11.49]
<i>Holdout: 2025, 416 assets, 249 days</i>		
Market only	0.00195 [0.00104, 0.00287]	−0.04 [−0.41, 0.34]
Market + assets	0.00502 [0.00412, 0.00593]	−1.49 [−1.87, −1.12]

References

Anderson, T.W., Darling, D.A., 1952. Asymptotic theory of certain “goodness of fit” criteria based on stochastic processes. The Annals of Mathematical

Statistics 23, 193–212.

- Bilmes, J.A., 1998. A Gentle Tutorial of the EM Algorithm and Its Application to Parameter Estimation for Gaussian Mixture and Hidden Markov Models. Technical Report TR-97-021. International Computer Science Institute.
- Booth, D.G., Fama, E.F., 1992. Diversification returns and asset contributions. *Financial Analysts Journal* 48, 26–32. doi:10.2469/faj.v48.n3.26.
- Bouchev, P., Nemtchinov, V., Wong, T.K.L., 2015. Volatility harvesting in theory and practice. *Journal of Wealth Management* 18, 89–100. doi:10.3905/jwm.2015.18.3.089.
- Bulla, J., Bulla, I., 2006. Stylized facts of financial time series and hidden semi-Markov models. *Computational Statistics & Data Analysis* 51, 2192–2209. doi:10.1016/j.csda.2006.07.021.
- Grinstead, C.M., Snell, J.L., 1997. Introduction to Probability. American Mathematical Society.
- Kolmogorov, A.N., 1933. Sulla determinazione empirica di una legge di distribuzione. *Giornale dell'Istituto Italiano degli Attuari* 4, 83–91.
- Kotz, S., Kozubowski, T.J., Podgórski, K., 2001. The Laplace Distribution and Generalizations: A Revisit with Applications to Communications, Economics, Engineering, and Finance. Birkhauser.
- Markowitz, H.M., 1976. Investment for the long run: New evidence for an old rule. *Journal of Finance* 31, 1273–1286. doi:10.1111/j.1540-6261.1976.tb03213.x.
- Rabiner, L., Juang, B., 1986. An introduction to hidden Markov models. *IEEE ASSP Magazine* 3, 4–16. doi:10.1109/MASSP.1986.1165342.
- Smirnov, N., 1948. Table for estimating the goodness of fit of empirical distributions. *The Annals of Mathematical Statistics* 19, 279–281.