

Hybrid Hidden Markov Model for Modeling Equity Excess Growth Rate Dynamics: A Discrete-State Approach with Jump-Diffusion

ABDULRAHMAN ALSWAIDAN, Cornell University, USA

JEFFREY VARNER, Cornell University, USA

Generating synthetic financial time series that preserve the statistical properties of real market data is essential for stress testing, risk model validation, and scenario design, yet existing approaches, from parametric models to deep generative networks, struggle to simultaneously reproduce the heavy-tailed distributions, negligible linear autocorrelation, and persistent volatility clustering that characterize empirical equity data. We propose a hybrid hidden Markov framework that discretizes continuous excess growth rates into Laplace quantile-defined market states and augments the regime-switching process with a Poisson-driven jump-duration mechanism to enforce realistic tail-state dwell times. Parameters are estimated by direct transition counting, avoiding the Baum-Welch EM algorithm entirely. We evaluate synthetic data quality using three complementary metrics: Kolmogorov-Smirnov and Anderson-Darling goodness-of-fit pass rates for distributional fidelity, and ACF mean absolute error for temporal structure preservation. Applied to ten years of SPY data across 1,000 simulated paths, the framework achieves KS and AD pass rates exceeding 97% and 90% in-sample and 95% out-of-sample (full calendar year 2025), reproducing the ARCH effect that standard regime-switching models miss. GARCH(1,1), by contrast, achieves only 4.6% and 1.1% in-sample pass rates despite its well-known volatility clustering properties. A Single-Index Model extension propagates the SPY factor path to a 424-asset universe, enabling scalable correlated synthetic path generation while preserving cross-sectional correlation structure.

Additional Key Words and Phrases: Synthetic Data Quality, Time Series Generation, Hidden Markov Model, Jump-Diffusion, Volatility Clustering, Regime-Switching, Stylized Facts, Distributional Fidelity

1 Introduction

The generation of synthetic financial time series that faithfully preserve the statistical properties of real market data is a central challenge in quantitative finance and data quality research. High-fidelity synthetic data are essential for stress testing risk models against plausible but unobserved market scenarios, validating portfolio optimization algorithms, and augmenting limited training sets for machine learning pipelines [2, 24]. Yet the quality requirements for financial synthetic data are unusually demanding: empirical equity excess growth rates exhibit three well-documented statistical regularities, namely heavy-tailed (leptokurtic) distributions, negligible linear autocorrelation in raw returns, and persistent volatility clustering, that any generative framework must simultaneously reproduce to be considered statistically faithful [4, 10, 30]. These domain-specific quality criteria, often called *stylized facts*, serve as the benchmark against which the fidelity of synthetic financial data must be evaluated.

Existing generative approaches address these quality requirements from different angles but each falls short on at least one dimension. GARCH-family models [8, 13] capture conditional variance dynamics but do not explicitly represent discrete market regimes or sudden jumps in volatility. Stochastic-volatility and jump-diffusion models [22, 32] enrich tail behavior but lack a mechanism for volatility persistence after jumps. Deep generative models, including recurrent neural networks [16, 23] and generative adversarial networks [29, 42, 43], can learn complex distributions from data but struggle to reproduce temporal dependence structures, particularly volatility clustering,

Authors' Contact Information: Abdulrahman Alswaidan, Cornell University, Robert Frederick Smith School of Chemical and Biomolecular Engineering, Ithaca, NY, USA; Jeffrey Varner, Cornell University, Robert Frederick Smith School of Chemical and Biomolecular Engineering, Ithaca, NY, USA, jdv27@cornell.edu.

without explicitly encoding the relevant inductive biases [4]. Hidden Markov models (HMMs) and related regime-switching frameworks offer a probabilistic structure for non-stationary market behavior through latent state dynamics [35], yet standard HMM specifications fail to generate persistent high-volatility regimes, leading to overly rapid reversion following extreme events and poor temporal quality in the synthetic output [9, 37]. Bridging this gap requires a generative framework that preserves both distributional fidelity and temporal structure while remaining computationally scalable.

We address this challenge by developing a hybrid HMM framework that discretizes continuous excess growth rates into quantile-based regimes and augments the resulting Markov process with a jump-duration mechanism that enforces volatility persistence, linking smooth diffusion to discontinuous shocks identified by jump-detection methods [3]. Using a 424-asset dataset with SPY as a case study, we estimate transition dynamics empirically, avoiding the EM algorithm entirely, and validate the model through in-sample and out-of-sample tests on distributional fit and volatility clustering. We further extend the univariate framework to a correlated multi-asset setting via a Single-Index Model (SIM) factor structure [39], enabling synthetic path generation across all 424 assets in a single generative pass.

This paper makes four contributions to synthetic financial data generation and quality assessment. First, we propose augmenting a discrete-state regime-switching model with a two-parameter Poisson jump-duration mechanism (ϵ, λ) that forces the model to dwell in high-volatility tail states for empirically realistic durations, providing a structural inductive bias that preserves the temporal quality (volatility clustering) of the synthetic output, addressing the key limitation identified in prior evaluations of both HMM-based [37] and GAN-based [29] financial data generators. Second, by partitioning excess growth rates into quantile-defined states via a Laplace CDF, we enable direct frequentist counting of regime transitions, eliminating the computational cost and initialization sensitivity of the Baum-Welch EM algorithm and making the framework scalable to a 424-asset synthetic data pipeline. Third, we conduct a rigorous quality evaluation using three complementary metrics (Kolmogorov-Smirnov and Anderson-Darling pass rates for distributional fidelity, and ACF mean absolute error for temporal structure preservation), all reported with standard errors. Using ten years of SPY data (2014–2024) for training and 249 trading days (full calendar year 2025) for out-of-sample testing, we demonstrate KS and AD pass rates of 97.1% (SE 0.5) and 90.4% (SE 0.9) in-sample and 95.0% (SE 0.7) and 95.4% (SE 0.7) out-of-sample across 1,000 simulated paths, with volatility clustering preserved in both windows. Finally, we extend the univariate framework to the full 424-asset universe through a Single-Index Model factor projection, enabling scalable correlated synthetic path generation while preserving cross-sectional correlation structure.

We find that a state resolution of $N = 100$ provides a balance between distributional fidelity and statistical robustness, and that the jump-duration extension is necessary for reproducing persistent high-volatility regimes that standard regime-switching models miss.

2 Related Work

Classical stylized facts and the failure of Gaussian diffusion

The quantitative study of financial market dynamics has been shaped by a tension between tractable theoretical frameworks and the empirical complexity of observed data. The Black-Scholes-Merton option pricing model [7] and the geometric Brownian motion (GBM) paradigm embedded within it assume that log-returns are independently and identically distributed Gaussian random variables. Under this assumption, asset prices follow a diffusion with constant drift and volatility, price paths are almost surely continuous, and the tails of the return distribution decay exponentially. These




Fig. 1. Empirical stylized facts of SPY daily excess growth rates (2014–2024). Panel (a) shows the leptokurtic marginal distribution; the Laplace fit substantially outperforms the Gaussian. Panel (b) confirms heavy tails via a normal Q-Q plot. Panels (c) and (d) contrast the near-zero autocorrelation of raw returns (consistent with the efficient markets hypothesis) against the persistent autocorrelation of absolute returns, motivating the jump-duration extension.

properties have well-known practical advantages: GBM admits closed-form option prices, portfolio variances aggregate linearly, and model parameters can be estimated with minimal data.

However, the Gaussian diffusion paradigm conflicts with a broad set of empirical regularities that were documented well before the widespread adoption of the model. Mandelbrot [30] demonstrated that the distributions of speculative price changes exhibit excess kurtosis; the probability of large deviations from the mean is far higher than Gaussian predictions, a phenomenon now described as *fat tails* or *leptokurtosis*. Fama [14] showed that, consistent with the efficient markets hypothesis, raw return series display negligible linear autocorrelation; daily or weekly returns are essentially unpredictable from their own past values. Yet Schwert [38] and others documented that the *magnitude* of returns, measured by absolute or squared returns, exhibits substantial positive autocorrelation that decays slowly over time. This temporal clustering of variance, where large price changes tend to be followed by large changes regardless of sign, is referred to as *volatility clustering* and is among the most persistently observed features of financial time series. Cont [10] provides a systematic review of these and related stylized facts, confirming their presence across asset classes, geographic markets, and sampling frequencies. Collectively, leptokurtosis, the absence

of return predictability, and volatility clustering define the benchmark against which any proposed model of market dynamics must be evaluated.

Conditional variance models: ARCH and GARCH families

The most influential class of models designed to address volatility clustering within the Gaussian framework is the autoregressive conditional heteroscedasticity (ARCH) family introduced by Engle [13]. The ARCH model conditions the variance of the current observation on the squared residuals of past observations, allowing variance to cluster without introducing autocorrelation in the level series. Bollerslev [8] generalized this to the GARCH(p, q) framework, in which the conditional variance depends on both past squared residuals and past conditional variances; the GARCH(1,1) specification became the canonical one-factor volatility model and remains widely used in practice. GARCH models are parsimonious, analytically tractable, and can be estimated by maximum likelihood under distributional assumptions on the innovation.

Despite their success in modeling conditional variance dynamics, standard GARCH models have several well-documented limitations in the context of stylized-facts reproduction. First, while GARCH innovations can be made leptokurtic by specifying Student- t or generalized error distributions, the fat-tail behavior arises through the parametric distributional assumption rather than through the structural dynamics of the model. Second, GARCH models do not explicitly represent discrete market regimes or the possibility of sudden jumps in the volatility process; the variance responds smoothly and continuously to past shocks. Third, regime persistence, the tendency of markets to remain in high-volatility states for sustained periods, is captured only indirectly through the persistence parameter of GARCH, which governs exponential decay of the volatility response. When persistence is near unity, the GARCH model is near-integrated and exhibits long-memory-like behavior, but this is a parametric approximation rather than a structural representation of distinct market regimes [12]. Stochastic volatility models [22, 41] address some of these shortcomings by treating variance as a latent diffusion process, but they require computationally intensive estimation via particle filters or Markov chain Monte Carlo methods, and their parameters do not directly correspond to observable market states.

Jump-diffusion models

An alternative approach to modeling fat tails and discontinuous price behavior is the explicit introduction of Poisson-driven jumps into the asset price process. Merton [32] extended the Black-Scholes framework by adding a compound Poisson jump component to the GBM, producing a model in which prices follow a diffusion between events and undergo discrete jumps at random times governed by a Poisson arrival process. This construction directly addresses leptokurtosis: the marginal distribution of returns is a Gaussian mixture weighted by the Poisson probability of observing zero, one, or more jumps over the sampling interval. Option pricing under Merton's model requires integrating over the jump count distribution, but retains analytical tractability through the characteristic function.

Kou [28] extended the jump-diffusion framework by replacing the Gaussian jump amplitude distribution with a double-exponential (asymmetric Laplace) distribution, which provides a better fit to the empirically observed asymmetry between the sizes of positive and negative jumps. This modification preserves analytical tractability for option pricing while improving the empirical fit to return distributions. Jump-diffusion models are particularly relevant in the context of equity markets, where identifiable macroeconomic announcements, earnings surprises, and geopolitical events generate discrete, large-magnitude price movements that are distinct from the continuous diffusion component [3]. A fundamental challenge, however, is that jump-diffusion models in their standard form do not provide an explicit mechanism for volatility persistence; the post-jump

volatility immediately returns to its baseline level, absent a separate stochastic volatility component. Integrating jump-diffusion with regime-switching provides a natural solution to this limitation, motivating the hybrid framework developed in the present study.

Hidden Markov models in finance

Hidden Markov models provide a probabilistic framework for non-stationary time series in which the observed data are generated by a process that switches among a discrete set of latent states according to Markov dynamics. The foundational algorithmic framework was developed by Rabiner and Juang [35], who introduced the Baum-Welch algorithm for maximum-likelihood estimation of the emission and transition parameters from observed sequences. Hamilton [19] applied the regime-switching HMM to macroeconomic time series, demonstrating that U.S. GDP growth is well described by a two-state model with distinct expansion and recession regimes; this work established the template for applying latent-state models to financial time series. The identification of time-varying market regimes has since become a central application of HMMs in finance, with applications ranging from stock return modeling to volatility forecasting and asset allocation.

Rydén, Teräsvirta, and Åsbrink [37] provided an influential systematic evaluation of HMM specifications against the stylized facts of daily return series. They showed that simple Gaussian-emission HMMs with a small number of states could reproduce leptokurtosis and the absence of return autocorrelation, but failed to generate the observed volatility clustering; ACF of absolute returns under standard HMMs decayed far too rapidly relative to empirical data. This diagnostic limitation identified the core structural gap that the present study addresses through the jump-duration extension. Bulla and Bulla [9] advanced this line of inquiry by considering hidden semi-Markov models, in which the dwell time in each state follows a state-specific distribution rather than the geometric distribution implied by standard Markov dynamics. They showed that the semi-Markov extension substantially improves the reproduction of volatility clustering by allowing the model to dwell in high-variance states for durations drawn from a flexible parametric family. Rossi and Gallo [36] applied HMMs to multivariate volatility estimation, demonstrating that latent state dynamics can capture the joint clustering of volatility across assets in a more interpretable manner than multivariate GARCH models. Nguyen [33] further demonstrated the utility of HMMs for stock trading, showing that regime identification from price data can inform tactical asset allocation.

A persistent limitation of standard HMM approaches in finance is the computational burden of maximum-likelihood estimation via the EM algorithm, particularly when the number of states is large or the asset universe is broad. The Baum-Welch procedure requires iterative forward-backward passes through the data and may converge to local maxima depending on initialization. For high-dimensional or large-state-space models, this creates a barrier to scalable application. The present study addresses this limitation directly by replacing EM with direct frequentist counting of transitions between quantile-defined states, an approach that is computationally trivial and free of initialization sensitivity.

Synthetic data generation and generative frameworks

The generation of synthetic financial time series has attracted increasing attention as both a practical tool (for augmenting limited data, testing risk models, and stress testing portfolios) and as a benchmark for evaluating the quality of generative models. Assefa et al. [2] provide a comprehensive survey of the opportunities, challenges, and pitfalls of synthetic data generation in finance, noting that stylized-facts reproduction is the primary criterion for evaluating the statistical quality of synthetic financial time series, and that naive generative approaches tend to fail on at least one of the three core stylized facts.

The rise of generative adversarial networks (GANs) prompted several investigations into whether deep learning architectures could learn stylized facts from financial data. Takahashi, Chen, and Tanaka-Ishii [42] showed that standard GAN architectures produce synthetic return sequences that visually resemble financial time series but fail to reproduce ACF structure precisely. Kwon and Lee [29] provided a careful evaluation of GAN-based methods against the stylized-facts checklist, finding that while GANs match marginal distributions reasonably well, capturing temporal dependence structures, particularly volatility clustering, remains challenging without explicitly encoding the relevant inductive biases into the architecture or training objective. Foundation models for Markov jump processes, as introduced by Berghaus et al. [5], offer a complementary perspective: large-scale pretrained inference networks can be applied to new Markov process instances without retraining, but their application to financial time series generation remains an open research direction.

Against this backdrop, the hybrid HMM framework of the present study occupies a distinct position: it is explicitly designed around the structural properties known to generate stylized facts, uses an interpretable discrete-state representation that can be audited and stress-tested, and achieves strong statistical fidelity with minimal computational overhead. The empirical estimation procedure avoids the training instability and data requirements of deep generative models, while the jump-duration mechanism provides a principled solution to the volatility-clustering limitation of standard HMMs. These properties make the framework particularly well-suited for deployment in risk management and stress-testing pipelines, where interpretability and computational reproducibility are as important as statistical fidelity.

Factor models and multi-asset extensions

Modeling the joint dynamics of large asset universes requires a framework for managing cross-sectional dependence. The Single-Index Model (SIM) of Sharpe [39] provides the canonical linear factor decomposition: each asset's excess return is decomposed into a common market factor component and an idiosyncratic residual, drastically reducing the number of parameters required to specify a covariance matrix relative to a full multivariate model. The SIM remains foundational in portfolio construction and risk management despite the availability of richer multi-factor models such as Fama-French [15], because its simplicity enables transparent attribution and scalable computation.

Existing HMM-based work in finance has been predominantly univariate, focusing on single-asset or single-index regime identification. Applications to individual equities [33], volatility modeling [36], and macro time series [19, 25] collectively demonstrate the utility of latent state representations, but leave open the question of how HMM-derived generative models can be extended to large cross-sections of assets in a computationally tractable way. Mixture hidden Markov models [11] offer a partial solution by allowing multiple assets to share a common latent state structure, but estimating such models jointly across hundreds of assets is computationally demanding.

The present study addresses this gap by combining the univariate HMM framework with the SIM factor decomposition. The hybrid HMM is estimated on the SPY index, which serves as the market-factor generator; asset-level paths are then reconstructed via the linear SIM projection. This construction exploits the fact that the dominant source of joint volatility clustering in equity markets is the common factor, and that idiosyncratic shocks can be treated as approximately independent given the factor. The result is a scalable generative pipeline for a 424-asset universe that inherits the stylized-facts fidelity of the HMM without requiring joint estimation of a high-dimensional model.

3 Methods

Data.

The empirical foundation of this study was a dataset composed of 424 United States-listed equities and exchange-traded funds spanning 10 years (2014-2024). A pipeline was developed to support the automated construction of excess growth rate models for any ticker within this dataset, allowing for batch simulation and stress testing across market sectors. To validate the performance of this pipeline, we focused our analysis on the single-asset SPY, which tracked the Standard & Poor's (S&P) 500 index. The data included daily open, high, low, and close prices along with volume-weighted average price metrics (VWAP). For training purposes, the first 2,766 observations constituted the in-sample dataset, spanning from January 3, 2014, to December 31, 2024. A subsequent 249 trading days, from January 2 to December 31, 2025, were reserved for out-of-sample testing to assess model generalizability.

Excess growth rate calculation.

We computed the excess growth rate $G_{i,j}$ for ticker i between time period $j - 1 \rightarrow j$ for a given equity price series $P_{i,j}$ (units: dollars per share (USD/share)). To isolate the stochastic premium earned over a risk-free asset, we computed the excess growth rate for ticker i during the period $j - 1 \rightarrow j$ as:

$$G_{i,j} \equiv \left(\frac{1}{\Delta t} \right) \cdot \ln \left(\frac{P_{i,j}}{P_{i,j-1}} \right) - r_f \quad (1)$$

where Δt represented the time step set to $1/252$ for daily data (units: years), and r_f denotes the constant continuously compounded risk-free rate (units: year^{-1}) derived from Separate Trading of Registered Interest and Principal of Securities (STRIPS) bond yields observed on September 10, 2025. Unlike return, the excess growth rate has units of inverse time (year^{-1}) and directly quantifies the premium earned above a continuously compounded risk-free benchmark. Log growth rates are time-additive and consistent with a continuous compounding framework. While r_f varied dynamically over the sample period, a constant proxy was employed here to isolate the endogenous volatility dynamics without introducing exogenous rate-curve noise.

Discrete hidden Markov model.

We defined the hidden Markov model using the tuple $\mathcal{M} = (\mathcal{S}, \mathcal{O}, \mathbf{T}, \mathbf{E}, \bar{\pi})$ [35]. The hidden state $S_t \in \mathcal{S} = \{1, \dots, N\}$ represented market mood regimes and was defined by quantile bins of a fitted cumulative distribution function (CDF) for the excess growth rate series. The observation space $\mathcal{O} \subset \mathbb{R}$ consisted of continuous excess growth rate values, the transition matrix $\mathbf{T}$ governed regime-to-regime dynamics, the emission model $\mathbf{E}$ was parameterized by state-conditional Normal distributions, and $\bar{\pi}$ denoted the stationary distribution. The observation model was

$$G_t \mid S_t = k \sim \mathcal{N}(\mu_k, \sigma_k). \quad (2)$$

To map observations to hidden states during estimation, we defined quantile boundaries $Q_k = F_L^{-1}(k/N; \mu_L, b_L)$ for $k \in \{1, \dots, N - 1\}$ and used finite outer bounds $Q_0 = F_L^{-1}(0.001; \mu_L, b_L)$ and $Q_N = F_L^{-1}(0.999; \mu_L, b_L)$. An observation was assigned to state k when $Q_{k-1} < R_t \leq Q_k$.

The choice of the Laplace distribution over alternatives such as the Student's t -distribution was motivated by its alignment with the peaks and heavy tails of high-frequency financial excess growth rates [10, 27]. While the Student's t -distribution was often used to model leptokurtosis, the Laplace distribution's peak at the mean matched the concentration of small price movements observed in the SPY dataset [21, 30]. Furthermore, the Laplace distribution provided a closed-form for quantile estimation, supporting the 424-ticker pipeline. While maximum likelihood estimation

(MLE) for the Student's t-distribution often required iterative numerical optimization, the Laplace distribution allowed for a direct MLE approach [6]. The quantile boundaries were derived from a parametric Laplace fit and therefore did not enforce equal counts per bin in the historical data; instead, they provided a partition of the distribution that preserved tail shape. By using the Laplace distribution to inform the binning process, the model used this baseline to characterize the peak and tails of the excess growth rate series before the empirical transition logic and jump-diffusion mechanism were applied [28].

We estimated the transition matrix T through direct frequentist counting rather than iterative optimization. This approach differed from the traditional expectation-maximization (EM) algorithm used for parameter estimation in Gaussian mixture and hidden Markov models [6]. However, this approach was justified because the discrete state assignments allowed for straightforward counting of observed transitions between regimes, and it was conceptually simple to implement and understand.

Empirical transition matrices often had low probabilities for exiting central states into tail states, which caused models to lose the effect of volatility clustering too quickly [9]. We corrected this by adding jump parameters ϵ , representing the probability of a jump event, and λ , representing the mean duration, along with a tail-state selection rule. Tail-states were defined as $S_{bottom} = \{1, \dots, N_{tail}\}$ and $S_{top} = \{N - N_{tail} + 1, \dots, N\}$, where N_{tail} was user-configurable. If a jump was triggered, the jump-duration K was sampled from a Poisson distribution such that $K \sim \text{Poisson}(\lambda)$, which was standard for modeling independent events in fixed intervals [17]. During these K steps, the standard Markovian transitions were overridden to target tail-states, with a user-configurable bias p_{neg} (default 0.52) that favored negative-tail states to reproduce gain/loss asymmetry.




Fig. 2. Architecture of the Hybrid Hidden Markov Model with Jump-Diffusion (HMM-WJ). At each step the chain transitions according to the empirical transition matrix T (probability $1 - \epsilon$, thin bidirectional arrows) or enters a Poisson jump episode (probability ϵ , dashed upward arrow to the Poisson clock). During a jump episode the active state is forced to the bottom tail set S_- (red, lowest-valued states) or the top tail set S_+ (blue, highest-valued states) for $K \sim \text{Poisson}(\lambda)$ consecutive steps before normal Markovian evolution resumes. Small bell curves above each state depict the state-conditional $N(\mu_k, \sigma_k)$ emission distribution. Tail sets are defined as the N_{tail} lowest- and highest-quantile states under the Laplace quantile partition.

To characterize the excess growth rates within each regime, we parameterized the emission model with Normal distributions $N(\mu_k, \sigma_k)$ whose bounds coincided with the 0.1 and 99.9 percentiles of the

state-conditional distribution. Specifically, we set $\mu_k = (Q_{k-1} + Q_k)/2$ and $\sigma_k = (Q_k - Q_{k-1})/(2z_{0.999})$, where $z_{0.999} \approx 3.09$ was the standard normal 99.9th percentile. During simulation, when the model was in state k , the continuous excess growth rate ($\hat{R}_t$) was sampled from the corresponding local distribution. This hybrid approach ensured that the model preserved the discontinuous components of the price process important for asset modeling while maintaining a smooth representation of intra-state variance [28].




Fig. 3. Fitted model internals for SPY with $N = 100$ states. Panel (a) overlays the fitted Laplace CDF on the empirical CDF; the 99 vertical lines mark the equal-probability quantile boundaries that define the N hidden states. The close agreement confirms that the Laplace distribution is an accurate marginal model for the excess growth rate data. Panel (b) displays the empirical transition matrix T in $\log_{10}$ colour scale; the dominant near-diagonal band reflects short-range regime persistence. Panel (c) plots the expected natural residence time $1/(1 - T_{kk})$ for each state on a $\log_{10}$ scale, with the bottom tail set S_- (states 1–5, red shading) and top tail set S_+ (states 96–100, teal shading) highlighted. Under pure Markovian dynamics every state, including the tail states, exits within 1–2 steps on average, far too quickly to reproduce empirical volatility clustering. The dashed horizontal line marks the mean jump duration $\lambda = 70$ steps imposed by the Poisson jump mechanism, quantifying the two-order-of-magnitude gap that motivates the jump-duration extension.

Growth rate generation.

Let's outline a simple sampling algorithm for generating random sequences from an HMM (without jumps). We initialize the generation process by drawing an initial hidden state X_0 from the stationary distribution $\bar{\pi}$. The stationary distribution $\bar{\pi}$ represented the long-run probability of the market occupying various regimes or moods [20]. While $\bar{\pi}$ mathematically satisfies the balance equation $\bar{\pi} = \bar{\pi}T$, we estimated this distribution numerically by raising the transition matrix T^n to a sufficiently large power $n = 50$, ensuring the vector reflected the equilibrium of market dynamics [18].

Given an HMM defined by the tuple $\mathcal{M} = (S, O, P, E, \bar{\pi})$, the number of time steps T , and the stationary distribution $\bar{\pi}$. Draw an initial hidden state X_0 from the initial distribution: $X_0 \sim \bar{\pi}$. For $n = 1$ to T , we sample the next hidden state X_n from the categorical distribution $\mathbb{P}_H(\dots)$ defined by the previous state X_{n-1} and the transition matrix T : $X_n \sim \mathbb{P}_H(X_n | X_{n-1})$. Next, we sample the observable state O_n from the categorical distribution $\mathbb{P}_E(\dots)$ defined by the current hidden state X_n and the emission matrix E : $O_n \sim \mathbb{P}_E(O_n | X_n)$. Finally, we store the sequences of hidden states and observable states. This procedure generates a sequences of hidden states $\{X_1, X_2, \dots, X_T\}$ and observable states $\{O_1, O_2, \dots, O_T\}$. The initial state X_0 is used to generate the sequence but is not included in the output.

Hyperparameter optimization via grid search.

To rigorously identify the optimal jump probability (ϵ) and mean jump duration (λ) [28, 32], we implemented a multi-objective grid search designed to minimize the discrepancy between the historical and simulated market signatures. The error function specifically targeted two stylized facts prevalent in financial time series [10, 37]: volatility clustering and leptokurtosis. We measured volatility clustering using the sum of squared errors between the observed and simulated autocorrelation function (ACF) of absolute excess growth rates [38] up to a lag of $L = 252$ trading days (approximately one trading year). To preserve the heavy-tailed nature of the distribution initially noted in foundational market studies [30], we included a penalty term for the global kurtosis.

The objective function $J(\epsilon, \lambda)$ was formalized as:

$$J(\epsilon, \lambda) = \sum_{\tau=1}^L \left(\text{ACF}_{obs}(\tau) - \overline{\text{ACF}}_{sim}(\tau) \right)^2 + w_K \left(K_{obs} - \bar{K}_{sim} \right)^2 \quad (3)$$

where $\text{ACF}_{obs}(\tau)$ and K_{obs} represented the empirical absolute growth autocorrelation at lag τ and the global kurtosis, respectively. The simulated equivalents, $\overline{\text{ACF}}_{sim}(\tau)$ and $\bar{K}_{sim}$, were averaged across 200 independent synthetic paths of 2,766 trading days per grid point [17]. A penalty weight of $w_K = 0.20$ was assigned to prioritize tail behavior alongside the temporal ACF decay. The grid search systematically explored ϵ values from 10^{-4} to 2.5×10^{-2} and λ values from 10 to 160. The full computational procedure is detailed in Algorithm 4.

Statistical validation metrics

To assess whether synthetic paths are statistically consistent with empirical data, we applied two nonparametric goodness-of-fit tests to the marginal distributions of simulated excess growth rate sequences. For each of the 1,000 simulated paths, the empirical cumulative distribution function (CDF) of the simulated sequence was compared against the empirical CDF of the historical in-sample or out-of-sample observations.

The Kolmogorov-Smirnov (KS) statistic is defined as:

$$D_{KS} = \sup_x |F_n(x) - F_m(x)| \quad (4)$$

where F_n and F_m denote the empirical CDFs of the observed and simulated sequences of lengths n and m , respectively [26, 40]. The null hypothesis of distributional equivalence is rejected at significance level α when D_{KS} exceeds the critical value $c(\alpha)\sqrt{(n+m)/(nm)}$.

The Anderson-Darling (AD) statistic places greater weight on the distributional tails than the KS test and is defined as:

$$A^2 = -n - \sum_{i=1}^n \frac{2i-1}{n} \left[\ln F(X_{(i)}) + \ln (1 - F(X_{(n+1-i)})) \right] \quad (5)$$

where $X_{(1)} \leq \dots \leq X_{(n)}$ are the order statistics of the pooled sample and F is the empirical CDF of the reference distribution [1]. The AD test is particularly suited for assessing tail-distributional fidelity, which is the primary concern for risk management applications.

We report the *pass rate* as the proportion of the 1,000 simulated paths for which the null hypothesis of distributional equivalence was not rejected at the $\alpha = 0.05$ significance level. A pass rate above 95 percent indicates that the synthetic paths are, in aggregate, statistically indistinguishable from the historical data under that test. Pass rates are reported separately for the in-sample and out-of-sample windows and for both model variants (HMM without jumps and the hybrid HMM with jumps).

Single-Index Factor Model extension

The framework described above applies to any individual asset ticker within the 424-asset dataset. To generalize synthetic path generation to a correlated multi-asset setting without fitting a separate HMM for each asset, we exploit the linear factor structure of the Single-Index Model (SIM) [39]. The SIM expresses the excess growth rate of asset i at time t as:

$$G_{i,t} = \alpha_i + \beta_i \cdot G_{\text{SPY},t} + \eta_{i,t}, \quad \eta_{i,t} \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, \sigma_{\eta,i}^2) \quad (6)$$

where $G_{\text{SPY},t}$ denotes the SPY excess growth rate at time t , β_i measures the systematic sensitivity of asset i to the index, α_i is the idiosyncratic drift, and $\eta_{i,t}$ is a zero-mean idiosyncratic shock assumed independent across assets and time. The parameters $(\alpha_i, \beta_i, \sigma_{\eta,i}^2)$ are estimated via ordinary least squares on the in-sample time series for each asset i .

Synthetic multi-asset paths are generated as follows. First, the hybrid HMM produces a single synthetic SPY excess growth rate path $\{\hat{G}_{\text{SPY},t}\}_{t=1}^T$ using the regime-switching and jump-duration mechanism described above. Second, for each asset i , an idiosyncratic shock sequence $\{\hat{\eta}_{i,t}\}$ is drawn independently from $\mathcal{N}(0, \hat{\sigma}_{\eta,i}^2)$. Third, the synthetic excess growth rate for asset i is recovered via:

$$\hat{G}_{i,t} = \hat{\alpha}_i + \hat{\beta}_i \cdot \hat{G}_{\text{SPY},t} + \hat{\eta}_{i,t}. \quad (7)$$

This construction preserves the cross-sectional correlation structure induced by the common factor while maintaining the regime-switching and jump-diffusion dynamics of the index. The resulting 424-dimensional synthetic path can be used for stress testing, portfolio risk assessment, and scenario generation across the full asset universe in a single generative pass, avoiding the curse of dimensionality that would accompany a full multivariate HMM.

Algorithms

The four algorithms below detail the complete computational pipeline: model construction (Algorithm 1), jump-diffusion simulation (Algorithm 2), state decoding and continuous rate reconstruction (Algorithm 3), and hyperparameter grid search (Algorithm 4).

Algorithm 1 Empirical hidden Markov model construction via Laplace partitioning

Require: Price history $P_{1:T}$, Risk-free rate r_f , Number of states N

Ensure: Transition matrix $\mathbf{T}$, Quantile boundaries $\mathbf{Q}$

- 1: Compute excess growth rates $R_t = \frac{1}{\Delta t} \ln(P_t/P_{t-1}) - r_f$ for all t .
 - 2: Fit Laplace parameters (μ_L, b_L) to the series $\mathbf{R}$ using maximum-likelihood estimation.
 - 3: Define interior boundaries $Q_k = F_L^{-1}(k/N; \mu_L, b_L)$ for $k \in \{1, \dots, N-1\}$, where F_L^{-1} is the inverse CDF of the Laplace distribution.
 - 4: Set finite outer bounds $Q_0 = F_L^{-1}(0.001; \mu_L, b_L)$ and $Q_N = F_L^{-1}(0.999; \mu_L, b_L)$.
 - 5: Assign each excess growth rate R_t to a discrete state $s_t \in \{1, \dots, N\}$ based on the boundaries in $\mathbf{Q} = \{Q_k\}$.
 - 6: Initialize count matrix $\mathbf{C} \in \mathbb{R}^{N \times N}$ and transition matrix $\mathbf{T} \in \mathbb{R}^{N \times N}$ with zeros.
 - 7: **for** $t = 1$ to $T-1$ **do**
 - 8: Identify transition $i = s_t$ to $j = s_{t+1}$.
 - 9: $\mathbf{C}_{i,j} \leftarrow \mathbf{C}_{i,j} + 1$.
 - 10: **end for**
 - 11: **for** $i = 1$ to N **do**
 - 12: $\mathbf{T}_{i,:} \leftarrow \mathbf{C}_{i,:} / \sum_{j=1}^N \mathbf{C}_{i,j}$ {Row-wise normalization}
 - 13: **end for**
 - 14: **return** $\mathbf{T}, \mathbf{Q}$
-

Algorithm 2 Hybrid jump-diffusion simulation with persistent volatility

Require: Model parameters $\mathbf{T}, \bar{\pi}, \epsilon, \lambda, N_{tail}, p_{neg}$, Total Steps M

Ensure: Simulated state sequence $S_{1:M}$

```

1: Define tail-states  $\mathcal{S}_{bottom} = \{1, \dots, N_{tail}\}$  and  $\mathcal{S}_{top} = \{N - N_{tail} + 1, \dots, N\}$ .
2: Sample initial state  $S_1 \sim \text{Categorical}(\bar{\pi})$ .
3: Set  $counter \leftarrow 2$ .
4: while  $counter \leq M$  do
5:   Sample  $u \sim \text{Uniform}(0, 1)$ .
6:   if  $u < \epsilon$  then
7:     Sample jump-duration  $K \sim \text{Poisson}(\lambda)$ .
8:     for  $j = 1$  to  $K$  while  $counter \leq M$  do
9:       Sample  $w \sim \text{Uniform}(0, 1)$ .
10:      if  $w < p_{neg}$  then
11:         $S_{counter} \sim \text{Uniform}(\mathcal{S}_{bottom})$  {Bias toward negative-tail events}
12:      else
13:         $S_{counter} \sim \text{Uniform}(\mathcal{S}_{top})$ 
14:      end if
15:       $counter \leftarrow counter + 1$ .
16:    end for
17:  else
18:    Sample  $S_{counter} \sim \text{Categorical}(\mathbf{T}_{S_{counter-1},:})$ .
19:     $counter \leftarrow counter + 1$ .
20:  end if
21: end while
22: return  $S_{1:M}$ 

```

Algorithm 3 State decoding and continuous excess growth rate reconstruction

Require: Simulated state sequence $S_{1:M}$, Quantile boundaries Q

Ensure: Reconstructed continuous excess growth rates $\hat{R}_{1:M}$

```

1: Define  $z_{0.999} = 3.09$ .
2: for  $k = 1$  to  $N$  do
3:   Set  $\mu_k = (Q_{k-1} + Q_k)/2$ .
4:   Set  $\sigma_k = (Q_k - Q_{k-1})/(2z_{0.999})$ .
5: end for
6: Initialize the reconstructed excess growth rate vector  $\hat{\mathbf{R}}$  of length  $M$ .
7: for  $t = 1$  to  $M$  do
8:   Retrieve the simulated state  $k = S_t$ .
9:   Sample the continuous excess growth rate  $\hat{R}_t$  from the normal distribution  $\mathcal{N}(\mu_k, \sigma_k)$ .
10: end for
11: return  $\hat{R}_{1:M}$ 

```

Algorithm 4 Multi-objective grid search for jump hyperparameters

Require: Observed growth rates $\mathbf{R}_{obs}$, ACF max lag $L = 252$, Number of paths $P = 1000$, Time steps $M = 2766$

Require: Grid space for $\epsilon \in [10^{-5}, 10^{-4}]$ and $\lambda \in [1.0, 100.0]$

Require: Kurtosis penalty weight $w_K = 0.20$

Ensure: Optimal jump hyperparameters (ϵ^*, λ^*)

```

1: Calculate observed absolute ACF:  $A_{obs}(\tau) = \text{ACF}(|\mathbf{R}_{obs}|, \tau)$  for  $\tau \in \{1, \dots, L\}$ .
2: Calculate observed global kurtosis:  $K_{obs} = \text{Kurtosis}(\mathbf{R}_{obs})$ .
3: Initialize minimum error  $E_{min} \leftarrow \infty$ .
4: for each  $\epsilon$  in grid space do
5:   for each  $\lambda$  in grid space do
6:     Initialize simulated ACF accumulator  $\mathbf{A}_{sum} \leftarrow \mathbf{0}$  and kurtosis accumulator  $K_{sum} \leftarrow 0$ .
7:     for  $p = 1$  to  $P$  do
8:       Simulate path using Algorithm 2 to obtain state sequence  $S_{1:M}$ .
9:       Decode sequence using Algorithm 3 to obtain continuous rates  $\hat{\mathbf{R}}^{(p)}$ .
10:       $\mathbf{A}_{sum}(\tau) \leftarrow \mathbf{A}_{sum}(\tau) + \text{ACF}(|\hat{\mathbf{R}}^{(p)}|, \tau)$  for all  $\tau$ .
11:       $K_{sum} \leftarrow K_{sum} + \text{Kurtosis}(\hat{\mathbf{R}}^{(p)})$ .
12:    end for
13:    Compute ensemble averages:  $\bar{A}_{sim}(\tau) = \mathbf{A}_{sum}(\tau)/P$  and  $\bar{K}_{sim} = K_{sum}/P$ .
14:    Compute objective:  $J = \sum_{\tau=1}^L \left( A_{obs}(\tau) - \bar{A}_{sim}(\tau) \right)^2 + w_K \left( K_{obs} - \bar{K}_{sim} \right)^2$ .
15:    if  $J < E_{min}$  then
16:       $E_{min} \leftarrow J$ 
17:       $(\epsilon^*, \lambda^*) \leftarrow (\epsilon, \lambda)$ 
18:    end if
19:  end for
20: end for
21: return  $\epsilon^*, \lambda^*$ 

```

4 Empirical Study

4.1 Descriptive statistics and stylized facts

Descriptive statistics for the SPY daily excess growth rate series confirm that the three canonical stylized facts are present in both the in-sample (2014–2024, $T = 2,766$) and out-of-sample (2025, $T = 249$) windows (Table 1). Excess kurtosis reaches 7.7 (IS) and 6.9 (OoS), the Jarque-Bera normality test is rejected ($p < 0.001$) in both windows, and the Ljung-Box test on $|G_t|$ is likewise rejected ($p < 0.001$), confirming leptokurtosis, negligible raw-return autocorrelation, and persistent volatility clustering. Any faithful generative model must reproduce all three properties simultaneously.

Table 1. Descriptive statistics and stylized-facts tests for SPY daily excess growth rates. In-sample: 2014–2024 ($T = 2,766$); out-of-sample: 2025 ($T = 249$). Gaussian and Laplace columns show MLE fits to the in-sample data. JB = Jarque-Bera normality test. LB = Ljung-Box autocorrelation test at lag 20. * denotes rejection at $\alpha = 0.05$.

Statistic	IS Observed	OoS Observed	Gaussian	Laplace
Mean (annualized, %)	6.31	11.60	6.31	14.19
Std Dev (annualized, %)	214.50	220.62	214.46	205.93
Skewness	−0.753	−1.093	0	0
Excess Kurtosis	7.715	6.867	0	3
JB normality test	reject* ($p < 0.001$)	reject* ($p < 0.001$)	n/a	n/a
LB test on G_t (lag 20)	reject* ($p < 0.001$)	reject* ($p = 0.019$)	n/a	n/a
LB test on $ G_t $ (lag 20)	reject* ($p < 0.001$)	reject* ($p < 0.001$)	n/a	n/a

4.2 Jump hyperparameter optimization and state resolution

A multi-objective grid search over $\epsilon \in [10^{-4}, 2.5 \times 10^{-2}]$ and $\lambda \in [10, 160]$, with 200 simulated paths per grid point, identified a well-defined interior minimum at $\epsilon^* = 10^{-4}$ and $\lambda^* = 100$, indicating that the jump hyperparameters are identifiable and that the objective landscape is not degenerate (Figure 4, panel a). At these optimal values, the ACF of $|G_t|$ closely tracks the observed SPY ACF across lags 1–252, confirming that the jump mechanism reproduces the empirical ARCH effect (Figure 4, panel b).




Fig. 4. Hyperparameter grid search over (ϵ, λ) for SPY ($N = 100, 200$ paths per grid point). Panel (a): heatmap of the objective function $J(\epsilon, \lambda)$ in $\log_{10}$ scale over the full search grid; the red star marks the optimal pair (ϵ^*, λ^*) . The well-defined interior minimum confirms that the jump hyperparameters are identifiable from data. Panel (b): ACF of $|G_t|$ at the optimal parameters computed from 500 jump-containing paths (10th–90th percentile band shaded; mean curve solid navy) compared with the SPY observed ACF (dashed red) and the 95% significance band (dotted black).

The primary analysis uses $N = 100$ states, but a sensitivity study over $N \in \{30, 60, 90, 100, 150, 200\}$ shows that the framework is robust to the choice of state resolution across a wide range (Table 2). KS and AD pass rates remain stable for both HMM-NJ ($\geq 98.9\%$ KS, $\geq 97.9\%$ AD) and HMM-WJ ($\geq 96.0\%$ KS, $\geq 88.6\%$ AD), while ACF-MAE for HMM-WJ improves slightly from 0.057 at $N = 30$ to 0.049 at $N = 200$. State-occupancy diagnostics confirm that every state at $N \leq 200$ contains at least 7 observations, ensuring adequate statistical support per state under the Laplace equal-probability partition. At $N = 350$, however, certain states are never visited in the historical record, leaving zero transition probabilities that break the ergodic structure and establishing a practical upper bound on state resolution.

Table 2. Sensitivity of model performance to state resolution N . All metrics computed over 1,000 simulated in-sample paths at $\alpha = 0.05$. ACF-MAE = mean absolute error of the ACF of $|G_t|$ over lags 1–252. Standard errors in parentheses: binomial SE for pass rates, bootstrap SE ($B = 500$) for ACF-MAE.

Model	N	KS pass (%)	AD pass (%)	ACF-MAE
HMM-NJ	200	99.6 (0.2)	98.5 (0.4)	0.059
HMM-NJ	150	99.6 (0.2)	98.4 (0.4)	0.059
HMM-NJ	100	99.0 (0.3)	98.8 (0.3)	0.059
HMM-NJ	90	99.4 (0.2)	98.1 (0.4)	0.059
HMM-NJ	60	99.4 (0.2)	98.1 (0.4)	0.059
HMM-NJ	30	98.9 (0.3)	97.9 (0.5)	0.059
HMM-WJ	200	96.0 (0.6)	88.6 (1.0)	0.049
HMM-WJ	150	96.9 (0.5)	88.8 (1.0)	0.050
HMM-WJ	100	96.8 (0.6)	91.7 (0.9)	0.053
HMM-WJ	90	97.4 (0.5)	91.0 (0.9)	0.052
HMM-WJ	60	96.8 (0.6)	92.4 (0.8)	0.054
HMM-WJ	30	98.0 (0.4)	96.8 (0.6)	0.057

4.3 In-sample distributional and temporal quality

Table 3 compares 1,000 simulated paths of 2,766 trading days from GARCH(1,1), HMM-NJ, and HMM-WJ using two-sample KS and AD tests at $\alpha = 0.05$. GARCH(1,1) achieves IS KS and AD pass rates of only 4.6% (SE 0.7) and 1.1% (SE 0.3), indicating that the GARCH marginal distribution is systematically inconsistent with the observed data despite its well-known volatility clustering properties. HMM-NJ achieves 99.4% (SE 0.2) and 98.5% (SE 0.4), and HMM-WJ achieves 97.1% (SE 0.5) and 90.4% (SE 0.9), confirming that the Laplace-mixture emission structure adequately captures the observed heavy-tailed distribution.

The temporal quality of the three models diverges sharply. GARCH achieves the lowest IS ACF-MAE (0.031) because its variance equation is explicitly designed to match autocorrelation structure, yet this advantage comes at the cost of distributional fidelity, a tradeoff that the quality metrics expose directly. HMM-NJ produces an ACF of $|G_t|$ that decays rapidly toward zero rather than exhibiting the slow empirical decay, confirming that the standard transition matrix alone is structurally incapable of generating persistent volatility clustering (Figure 5, panel b). HMM-WJ, with an IS ACF-MAE of 0.052, generates a mixture of two path families: approximately 76% contain no Poisson jump and behave identically to HMM-NJ, while the remaining 24% contain at least one jump event and sustain substantial ACF($|G_t|$) decay across all lags. The jump-duration mechanism is therefore the structural feature responsible for reproducing the ARCH effect.

Table 3. In-sample and out-of-sample model comparison for SPY ($N = 100$, 1,000 simulated paths, significance level $\alpha = 0.05$). GARCH: GARCH(1,1) estimated by quasi-maximum likelihood. HMM-NJ: standard Hidden Markov Model without jumps. HMM-WJ: hybrid model with Poisson jump-duration extension. ACF-MAE: mean absolute error of the ACF of $|G_t|$ over lags 1–252. Kurtosis: mean excess kurtosis across simulated paths. Standard errors in parentheses: binomial SE for pass rates, $\text{std}/\sqrt{n}$ for kurtosis, bootstrap SE ($B = 500$) for ACF-MAE.

Metric	GARCH(1,1)	HMM-NJ ($N = 100$)	HMM-WJ ($N = 100$)
<i>In-sample: 2,766 trading days (2014–2024)</i>			
KS pass rate (%)	4.6 (0.7)	99.4 (0.2)	97.1 (0.5)
AD pass rate (%)	1.1 (0.3)	98.5 (0.4)	90.4 (0.9)
Excess kurtosis (observed)		7.715	
Excess kurtosis (simulated)	7.5 (0.32)	5.7 (0.02)	5.5 (0.03)
ACF-MAE (lags 1–252)	0.031	0.059	0.052
<i>Out-of-sample: 249 trading days (2025)</i>			
KS pass rate (%)	79.0 (1.3)	96.7 (0.6)	95.0 (0.7)
AD pass rate (%)	71.2 (1.4)	96.3 (0.6)	95.4 (0.7)
Excess kurtosis (observed)		6.867	
Excess kurtosis (simulated)	1.7 (0.07)	5.1 (0.07)	5.0 (0.08)
ACF-MAE (lags 1–252)	0.026	0.041	0.039
Parameters estimated	3	2	4

4.4 Out-of-sample evaluation

Out-of-sample testing over 249 trading days (January 2 through December 31, 2025) reveals that the HMM variants generalise robustly while GARCH degrades substantially (Table 3). GARCH achieves OoS KS and AD pass rates of 79.0% (SE 1.3) and 71.2% (SE 1.4); HMM-NJ achieves 96.7% (SE 0.6) and 96.3% (SE 0.6); HMM-WJ achieves 95.0% (SE 0.7) and 95.4% (SE 0.7). The simulated kurtosis for GARCH collapses to 1.7 against an observed value of 6.9, while the HMM variants maintain simulated kurtosis near 5.0–5.1, indicating that GARCH fails to preserve tail behaviour outside the training window.

Nonetheless, the out-of-sample period exposes the stationarity assumption embedded in the IS-calibrated jump parameters. The distribution of KS p -values shifts leftward from IS to OoS, reflecting the expected distributional degradation when stationary parameters are applied to a period with elevated macro uncertainty (Figure 6, panels a–b). The observed ACF of $|G_t|$ in the 2025 window exceeds the HMM-WJ simulation band at medium-to-long lags, indicating that (ϵ^*, λ^*) underestimate the frequency and persistence of volatility episodes in this period and that the jump parameters are regime-dependent (Figure 6, panel d). Despite this ACF gap, the observed marginal density remains within the simulation envelope, confirming adequate distributional coverage even under regime shift (Figure 6, panel c).




Fig. 5. Head-to-head in-sample model comparison for SPY ($N = 100, 1,000$ simulated paths). Panel (a): marginal density of excess growth rates with IS KS pass rates annotated. GARCH(1,1) fails the two-sample KS test on 94% of paths (pass rate 6%), indicating a systematic shape mismatch with the empirical distribution. Both HMM variants pass at rates $\geq 96\%$, confirming that the Laplace-mixture emission adequately captures the observed heavy tails. HMM-WJ generalises robustly out-of-sample (OoS KS pass rate 95%), while GARCH recovers partially (OoS 78%). Panel (b): autocorrelation function of absolute excess growth rates, $\text{ACF}(|G_t|)$, at lags 1–252 (shaded bands: 10th–90th percentile across paths). HMM-NJ is structurally incapable of producing persistent volatility clustering: without a jump mechanism the latent states are i.i.d. conditional on the Markov chain, so $\text{ACF}(|G_t|) \approx 0$ for all lags beyond one (dotted curve). HMM-WJ generates a *mixture* of two path families under the same model: approximately 76% of IS paths contain no Poisson jump and therefore behave identically to HMM-NJ (dashed curve, near zero); the remaining $\sim 24\%$ of paths contain at least one jump event and sustain substantial $\text{ACF}(|G_t|)$ decay across all lags (solid navy curve $\pm$ band). The jump frequency, and hence the fraction of the ensemble exhibiting slow ACF decay, is controlled by the hyperparameter λ , making the volatility-clustering strength directly tunable. Panel (c): tail Q-Q plot at the 0.1st–99.9th percentile region; HMM-WJ provides the closest mean quantile match in the extreme tails.




Fig. 6. Statistical validation of HMM-WJ ($N = 100, 1,000$ simulated paths, $\alpha = 0.05$). Panels (a) and (b) show the distribution of two-sample KS p -values in-sample and out-of-sample, respectively; a well-calibrated generative model produces p -values that are approximately uniform above the significance threshold. The leftward shift from (a) to (b) reflects the expected distributional degradation when stationary parameters are applied to a regime with elevated macro uncertainty. Panel (c) shows the out-of-sample marginal density fan chart; the observed density (dashed red) falls within the 10th–90th percentile simulation envelope, confirming adequate distributional coverage despite the regime shift. Panel (d) shows the ACF of $|G_t|$ in the out-of-sample window. The observed ACF (dashed red) exceeds the simulation band at medium-to-long lags, indicating that the 2025 test period exhibited stronger volatility clustering than the IS-calibrated jump parameters (ϵ^* , λ^*) predict. This gap provides direct evidence that the jump hyperparameters are regime-dependent and motivates the time-varying extensions discussed in Section 5.

4.5 Multi-asset extension

Propagating HMM-WJ SPY paths through a Single-Index Model for 424 S&P 500 constituents, with $\hat{G}_{i,t} = \hat{\alpha}_i + \hat{\beta}_i G_{\text{SPY},t} + \hat{\eta}_{i,t}$ estimated by OLS on the training window and idiosyncratic shocks drawn by resampling empirical residuals, yields a median IS KS pass rate of 66.7% and a mean of 58.4% across the universe (Figure 7, panel b; Table 4). These rates are lower than the 97.1% achieved for SPY alone and reflect the well-known limitation that a single-factor linear decomposition with symmetric residuals cannot fully capture asset-specific tail behaviour, skewness, or sector-level dynamics. The SIM R_i^2 varies widely across GICS sectors, from near-zero for defensive and low-correlation names to above 0.8 for broad-market ETFs (Figure 7, panel a), and the degradation in KS pass rate concentrates among assets with low R_i^2 . High- β assets closely tracking SPY retain pass rates above 80%, suggesting that the HMM-WJ generative engine is not the bottleneck; the limiting factor is the expressiveness of the factor model (Figure 7, panel c).

Out-of-sample KS pass rates across 417 surviving assets reach a mean of 82.1% and a median of 91.8%, confirming that the SIM extension generalises to the test window (Table 4). The correlation-distance distribution between simulated and historical paths yields a mean distance of $1 - \rho \approx 1.0$, confirming that the framework produces temporally independent realisations rather than tracking specific historical events, consistent with its role as a scenario generation tool rather than a forecasting instrument (Figure 6).

Table 4. Cross-sectional summary statistics for the Single-Index Model (SIM) extension across S&P 500 constituents. Individual asset paths are generated as $\hat{G}_{i,t} = \hat{\alpha}_i + \hat{\beta}_i G_{\text{SPY},t} + \hat{\eta}_{i,t}$ where $G_{\text{SPY},t}$ is drawn from HMM-WJ ($N = 100$) simulated paths and $\hat{\eta}_{i,t}$ is resampled from empirical residuals. KS pass rates computed over 1,000 paths at $\alpha = 0.05$.

Statistic	Mean	Median	5th pct	95th pct
<i>In-sample (424 assets, 2,766 trading days, 2014–2024)</i>				
KS pass rate (%)	58.4	66.7	2.5	96.0
SIM R^2	0.298	0.295	0.107	0.507
$\hat{\beta}$ (factor loading)	1.066	1.086	0.480	1.657
$\hat{\alpha}$ (intercept)	−0.0328	−0.0196	−0.2037	0.0964
<i>Out-of-sample (417 assets, 249 trading days, 2025)</i>				
KS pass rate (%)	82.1	91.8	23.8	99.0
SIM R^2	0.142	0.142	−0.169	0.479
$\hat{\beta}$ (factor loading)	1.064	1.084	0.482	1.660
$\hat{\alpha}$ (intercept)	−0.0320	−0.0186	−0.2035	0.0965




Fig. 7. Multi-asset extension via the Single-Index Model (SIM) across 424 S&P 500 constituents. Panel (a) summarises the SIM regression fit quality by GICS sector. Panel (b) shows that KS distributional consistency is maintained across the full asset universe when paths are generated via the factor structure $\hat{G}_{i,t} = \hat{\alpha}_i + \hat{\beta}_i G_{SPY,t} + \hat{\eta}_{i,t}$. Panel (c) illustrates representative marginal density comparisons for three assets spanning a wide range of systematic risk exposure.

5 Discussion

The empirical results establish that the hybrid HMM framework generates synthetic equity paths whose distributional and temporal properties closely match those of historical data, both in-sample and out-of-sample. In this section we interpret the key findings, discuss their practical implications, and identify the limitations that inform the scope of these conclusions.

5.1 The role of the jump-duration mechanism

The comparison between HMM-NJ and HMM-WJ isolates the contribution of the Poisson jump-duration mechanism. Without jumps, the Markov chain transitions out of tail states within one to two steps on average, producing an ACF of $|G_t|$ that decays too quickly and failing to reproduce the volatility clustering observed in the data. Introducing jumps forces the model to dwell in high-volatility regimes for empirically realistic durations, and approximately 24% of simulated paths contain at least one such episode, sufficient to shift the ensemble-level ACF into close agreement with the empirical profile. This result is consistent with the broader literature on regime-switching models, where standard HMM specifications have repeatedly been found to undergenerate volatility persistence [9, 37], and it demonstrates that a minimal two-parameter extension (ϵ, λ) can resolve the issue without the computational overhead of hierarchical or semi-Markov alternatives.

The tradeoff between distributional fidelity and temporal structure, visible in the GARCH comparison, deserves emphasis. GARCH(1,1) achieves the lowest ACF-MAE because its variance equation is explicitly parameterised to match autocorrelation dynamics, yet its marginal distribution is systematically inconsistent with the data (IS KS pass rate of 4.6%). The hybrid HMM reverses this tradeoff: the Laplace-mixture emissions deliver high distributional fidelity, and the jump mechanism provides sufficient temporal structure. Neither model dominates on all metrics simultaneously, but the HMM framework provides a more balanced quality profile for synthetic data applications where both distributional and temporal fidelity matter.

5.2 Practical implications

The proposed framework offers practical utility for both risk management and the broader synthetic data community. Generative models that reproduce the stylized facts of financial markets enable stress testing under market regimes that are statistically plausible but not historically observed, addressing a key limitation of scenario libraries derived solely from empirical records. The interpretable regime structure, where each hidden state corresponds to a quantile-defined segment of the excess growth rate distribution, facilitates communication between quantitative analysts and risk managers; states can be labeled (e.g., *crash*, *bear*, *neutral*, *bull*, *rally*) and linked to economic narratives. The model's computational efficiency further supports Monte Carlo-based risk metrics such as Value-at-Risk and Conditional Value-at-Risk at scale across large asset universes [17].

The avoidance of the Baum-Welch EM algorithm, made possible by the quantile-based state partition that explicitly assigns each observation to a hidden state, eliminates convergence sensitivity to initialisation and reduces computational cost. This property matters for large-scale pipelines: fitting the model to 424 assets requires only empirical counting and a two-parameter grid search, making it feasible to regenerate synthetic scenarios on a daily or weekly cadence as new data arrive [31, 34].

More broadly, the quality evaluation methodology demonstrated here, combining distributional fidelity tests (KS, AD) with temporal structure metrics (ACF-MAE) and reporting standard errors across Monte Carlo ensembles, provides a reusable template for assessing synthetic time series quality in other domains where temporal dependence structures are critical to downstream applications [4].

5.3 Limitations

Several limitations inform the interpretation of these results. First, the model assumes stationarity of the transition matrix and jump hyperparameters across the full in-sample window. The out-of-sample ACF analysis (Figure 6, panel d) provides direct evidence of this limitation: the IS-calibrated (ϵ^*, λ^*) underestimate volatility clustering persistence in the 2025 test window, consistent with the elevated macro uncertainty of that period. Structural breaks in market microstructure, such as those accompanying regulatory changes or the COVID-19 shock, may not be adequately captured by a single static transition matrix.

Second, the quantile-based state partition is fixed at estimation time and does not adapt to evolving volatility regimes. The boundaries are calibrated on the full in-sample distribution and may overweight the central mass relative to the tails during periods of prolonged market stress.

Third, the Single-Index Model extension imposes a linear, single-factor model of cross-sectional dependence. The median IS KS pass rate of 66.7% across 424 assets confirms that this decomposition cannot capture the heterogeneous tail behaviour, skewness, and sector-specific dynamics present in individual equities [15, 39]. However, the architecture is modular: the HMM-WJ generative engine for the market factor is decoupled from the factor model used to propagate paths, and replacing the SIM with a richer specification would not require modifying the underlying Markov chain machinery.

Finally, the out-of-sample evaluation covers a single 249-day window (full calendar year 2025), which limits the statistical power of out-of-sample inference. A longer or rolling evaluation would provide a more robust assessment of generalisation performance.

6 Conclusion

We developed a hybrid hidden Markov model that discretizes equity excess growth rates into Laplace quantile-defined states and augments regime-switching dynamics with a Poisson-driven jump-duration mechanism. Applied to SPY over a ten-year training window and validated out-of-sample on the full calendar year 2025, the framework achieves KS and AD pass rates exceeding 95% while reproducing the volatility clustering that standard regime-switching models miss. The jump mechanism, governed by only two scalar hyperparameters, is the structural feature responsible for generating persistent high-volatility regimes, and the direct frequentist estimation strategy eliminates the computational cost and initialisation sensitivity of the EM algorithm. A Single-Index Model extension scales the approach to a 424-asset universe, providing a baseline for multi-asset synthetic scenario generation.

Several extensions follow naturally. Time-varying transition matrices, estimated via rolling windows or Bayesian online learning, would allow the model to adapt to structural regime shifts. Replacing the single-factor SIM with multi-factor specifications such as the Fama-French three- or five-factor model [15], principal-component-based factor structures, or nonlinear residual generators could substantially improve asset-level distributional fidelity. Adaptive state definitions based on density clustering rather than fixed quantiles could improve tail coverage during crisis periods. Finally, embedding the hybrid HMM as the generative component within a portfolio optimisation loop, where synthetic scenarios drive tail-risk objectives, would provide an end-to-end application of the framework.

Conflict of Interest Statement

The authors declare that the research was conducted without any commercial or financial relationships that could potentially create a conflict of interest.

Author Contributions

J.V. directed the study. A.A. developed the model and simulation code, conducted the in-sample and out-of-sample analysis, and generated the figures. Both authors edited and reviewed the final manuscript.

Data Availability Statement

The model code, simulation scripts, training and testing data are available under a Massachusetts Institute of Technology (MIT) license from the GitHub repository: <https://github.com/varnerlab/HMM-Model-with-Jumps.git>

References

- [1] Theodore W. Anderson and Donald A. Darling. 1952. Asymptotic theory of certain “goodness of fit” criteria based on stochastic processes. *The Annals of Mathematical Statistics* 23, 2 (1952), 193–212.
- [2] Samuel A. Assefa, Danial Dervovic, Mahmoud Mahfouz, Robert E. Tillman, Prashant Reddy, and Manuela Veloso. 2021. Generating synthetic data in finance: opportunities, challenges and pitfalls. In *Proceedings of the First ACM International Conference on AI in Finance*. 1–8. doi:10.1145/3383455.3422554
- [3] Jay F. K. Au Yeung, Zi-kai Wei, Kit Yan Chan, Henry Y. K. Lau, and Ka-Fai Cedric Yiu. 2020. Jump detection in financial time series using machine learning algorithms. *Soft Computing* 24, 3 (2020), 1789–1801. doi:10.1007/s00500-019-04006-2
- [4] Kevin Bauer, Simon Joussen, and Jan Lukas. 2024. Thinking in Categories: A Survey on Assessing the Quality for Time Series Synthesis. *Journal of Data and Information Quality* (2024). doi:10.1145/3666006
- [5] David Berghaus, Kostadin Cvejovski, Patrick Seifner, César Ojeda, and Ramsés J. Sanchez. 2025. Foundation inference models for Markov jump processes. In *Proceedings of the 38th International Conference on Neural Information Processing Systems*. 129407–129442.
- [6] Jeff A. Bilmes. 1998. A Gentle Tutorial of the EM Algorithm and its Application to Parameter Estimation for Gaussian Mixture and Hidden Markov Models. *International Computer Science Institute* (1998).
- [7] Fischer Black and Myron Scholes. 1973. The Pricing of Options and Corporate Liabilities. *Journal of Political Economy* 81, 3 (1973), 637–654. doi:10.1086/260062
- [8] Tim Bollerslev. 1986. Generalized Autoregressive Conditional Heteroscedasticity. *Journal of Econometrics* 31, 3 (1986), 307–327. doi:10.1016/0304-4076(86)90063-1
- [9] Jan Bulla and Ingo Bulla. 2006. Stylized facts of financial time series and hidden semi-Markov models. *Computational Statistics & Data Analysis* 51, 4 (2006), 2192–2209. doi:10.1016/j.csda.2006.07.021
- [10] R. Cont. 2001. Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance* 1, 2 (2001), 223–236. doi:10.1080/713665670
- [11] José G. Dias, Jeroen K. Vermunt, and Sofia Ramos. 2010. Mixture Hidden Markov Models in Finance Research. In *Advances in Data Analysis, Data Handling and Business Intelligence*. Springer, 451–459. doi:10.1007/978-3-642-01044-6_41
- [12] Francis X. Diebold and Atsushi Inoue. 2001. Long memory and regime switching. *Journal of Econometrics* 105, 1–2 (2001), 131–159. doi:10.1016/S0304-4076(01)00073-2
- [13] Robert F. Engle. 1982. Autoregressive Conditional Heteroscedasticity with Estimates of the Variance of United Kingdom Inflation. *Econometrica* 50, 4 (1982), 987–1007.
- [14] Eugene F. Fama. 1970. Efficient Capital Markets: A Review of Theory and Empirical Work. *The Journal of Finance* 25, 2 (1970), 383–417. doi:10.2307/2325486
- [15] Eugene F. Fama and Kenneth R. French. 1993. Common risk factors in the returns on stocks and bonds. *Journal of Financial Economics* 33, 1 (1993), 3–56. doi:10.1016/0304-405X(93)90023-5
- [16] Thomas Fischer and Christopher Krauss. 2018. Deep learning with long short-term memory networks for financial market predictions. *European Journal of Operational Research* 270, 2 (2018), 654–669. doi:10.1016/j.ejor.2017.11.054
- [17] Paul Glasserman. 2003. *Monte Carlo Methods in Financial Engineering*. Springer. doi:10.1007/978-0-387-21617-1
- [18] Charles M. Grinstead and J. Laurie Snell. 1997. *Introduction to Probability*. American Mathematical Society.
- [19] James D. Hamilton. 1989. A New Approach to the Economic Analysis of Nonstationary Time Series and the Business Cycle. *Econometrica* 57, 2 (1989), 357–384. doi:10.2307/1912559
- [20] James D. Hamilton. 2018. Regime Switching Models. *The New Palgrave Dictionary of Economics* (2018), 11421–11426. doi:10.1057/978-1-349-95189-5_2459
- [21] V. Harwood. 2019. Against the Norm: Modeling Daily Stock Returns with the Laplace Distribution. *arXiv preprint arXiv:1906.10325* (2019).
- [22] Steven L. Heston. 1993. A Closed-Form Solution for Options with Stochastic Volatility with Applications to Bond and Currency Options. *The Review of Financial Studies* 6, 2 (1993), 327–343. doi:10.1093/rfs/6.2.327
- [23] Sepp Hochreiter and Jürgen Schmidhuber. 1997. Long Short-Term Memory. *Neural Computation* 9, 8 (1997), 1735–1780. doi:10.1162/neco.1997.9.8.1735
- [24] James Jordon, Lukasz Szpruch, Florimond Houssiau, Luke Sheridan Rains, James Sheridan, Michael Sheridan, and David Sheridan. 2022. Synthetic Data – what, why and how? *arXiv preprint arXiv:2205.03257* (2022).
- [25] Chang-Jin Kim and Charles R. Nelson. 1999. Has the U.S. Economy Become More Stable? A Bayesian Approach Based on a Markov-Switching Model of the Business Cycle. *The Review of Economics and Statistics* 81, 4 (1999), 608–616. doi:10.1162/003465399558472
- [26] A. N. Kolmogorov. 1933. Sulla determinazione empirica di una legge di distribuzione. *Giornale dell’Istituto Italiano degli Attuari* 4 (1933), 83–91.
- [27] Samuel Kotz, Tomasz J. Kozubowski, and Krzysztof Podgórski. 2001. *The Laplace Distribution and Generalizations: A Revisit with Applications to Communications, Economics, Engineering, and Finance*. Birkhauser.

- [28] S. G. Kou. 2002. A Jump-Diffusion Model for Option Pricing. *Management Science* 48, 8 (2002), 1086–1101. doi:10.1287/mnsc.48.8.1086.166
- [29] Sohyeon Kwon and Yongjae Lee. 2024. Can GANs Learn the Stylized Facts of Financial Time Series?. In *Proceedings of the 5th ACM International Conference on AI in Finance*. 126–133. doi:10.1145/3677052.3698661
- [30] Benoit Mandelbrot. 1963. The Variation of Certain Speculative Prices. *The Journal of Business* 36, 4 (1963), 394–419.
- [31] Simone Manganelli and Robert F. Engle. 2004. Value at risk models in finance. *European Central Bank Working Paper Series* (2004).
- [32] Robert C. Merton. 1976. Option Pricing When Underlying Stock Returns Are Discontinuous. *Journal of Financial Economics* 3, 1-2 (1976), 125–144. doi:10.1016/0304-405X(76)90022-2
- [33] Nguyet Nguyen. 2018. Hidden Markov Model for Stock Trading. *International Journal of Financial Studies* 6, 2 (2018), 36. doi:10.3390/ijfs6020036
- [34] Maureen O’Hara. 1998. *Market Microstructure Theory*. Wiley.
- [35] L. Rabiner and B. Juang. 1986. An introduction to hidden Markov models. *IEEE ASSP Magazine* 3, 1 (1986), 4–16. doi:10.1109/MASSP.1986.1165342
- [36] Alessandro Rossi and Giampiero M. Gallo. 2006. Volatility estimation via hidden Markov models. *Journal of Empirical Finance* 13, 2 (2006), 203–230. doi:10.1016/j.jempfin.2005.09.003
- [37] Tobias Rydén, Timo Teräsvirta, and Stefan Åsbrink. 1998. Stylized facts of daily return series and the hidden Markov model. *Journal of Applied Econometrics* 13, 3 (1998), 217–244. doi:10.1002/(SICI)1099-1255(199805/06)13:3<217::AID-JAE476>3.0.CO;2-V
- [38] G. William Schwert. 1989. Why Does Stock Market Volatility Change Over Time? *The Journal of Finance* 44, 5 (1989), 1115–1153. doi:10.1111/j.1540-6261.1989.tb02647.x
- [39] William F. Sharpe. 1963. A Simplified Model for Portfolio Analysis. *Management Science* 9, 2 (1963), 277–293. doi:10.1287/mnsc.9.2.277
- [40] N. Smirnov. 1948. Table for estimating the goodness of fit of empirical distributions. *The Annals of Mathematical Statistics* 19, 2 (1948), 279–281.
- [41] Elias M. Stein and Jeremy C. Stein. 1991. Stock price distributions with stochastic volatility: an analytic approach. *The Review of Financial Studies* 4, 4 (1991), 727–752.
- [42] Shuntaro Takahashi, Yu Chen, and Kumiko Tanaka-Ishii. 2019. Modeling financial time-series with generative adversarial networks. *Physica A: Statistical Mechanics and its Applications* 527 (2019), 121261. doi:10.1016/j.physa.2019.121261
- [43] Jinsung Yoon, Daniel Jarrett, and Mihaela van der Schaar. 2019. Time-series Generative Adversarial Networks. In *Advances in Neural Information Processing Systems*, Vol. 32.

Supplemental Information